

Mar 20, 06 16:52

ddr.photo

Page 1/8

```

/*****
* lab2.cpp
*
* Name: Mark Randles
* Class: CS408
* Date: 2006-02-28
*
* Description: Real-time simulation of disk
* access, based on a total service time for
* 1000 disk access requests.
*****/

/*****
* INCLUDES
*****/
#include <iostream>
using std::cout;
using std::cerr;
using std::endl;

#include <map>
using std::map;

#include <unistd.h>
#include <pthread.h>

#include "ddr.h"

/*****
* GLOBALS
*****/
bool done = false; // set to true if the disk drive has completed REQUESTS requests
RequestQueue *request_queue = NULL; // the request queue for the disk drive
pthread_mutex_t mutex_request_queue = PTHREAD_MUTEX_INITIALIZER;
bool complete[N]; // set to true when the disk_drive has completed the last request

/*****
* PROTOTYPES
*****/
void *disk_drive();
void *request_thread(void*);

/*****
* MAIN FUNCTION
*****/
int main(int argc, char* argv[]) {
    pthread_t threads[N]; // index of threads

    request_queue = NULL;

    cout << argc << endl;

    if(argc > 1 && (strcmp(argv[1], "--debug") == 0 || strcmp(argv[1], "-d") == 0)) {
        debug = true;
        DEBUG_PRINT("Debug ON" << endl);
        if(argc >= 3) {
            if(strcmp(argv[2], "--scan") == 0)
                request_queue = new RequestQueueSCAN();
        } else
    }

```

Monday March 20, 2006

Mar 20, 06 16:52

ddr.photo

Page

```

        request_queue = new RequestQueueFIFO();
    } else if(argc > 1 && (strcmp(argv[1], "--scan") == 0))
        request_queue = new RequestQueueSCAN();
    else
        request_queue = new RequestQueueFIFO();

    // make sure a request queue exists
    if(request_queue == NULL) {
        cerr << "No request queue created!";
        return(1);
    }

    // seed the random # generator
    srand(0);

    // create the request threads
    for(int i=0; i < N; i++)
        pthread_create(&threads[i], NULL, request_thread, (void *)i);

    // run the disk_drive() process, so we've only got N+1 threads
    disk_drive();

    // wait for all the threads to rejoin
    for(int i=0; i < N; i++)
        pthread_join(threads[i], NULL);

    // print some stats for the queue
    request_queue->print_stats();

    return(0);
}

/*****
* FUNCTIONS
*****/
void *disk_drive() {
    double total_time = 0.0; // sum of all the times
    int served = 0; // the total number of request processed, should = N

    double seek_time = 0.0; // the seek time for the request
    request *r = NULL;

    // do a service loop until we've processed N requests
    while(served < REQUESTS) {
        // see if there's a waiting request
        while(request_queue->request_count() <= 0) {
            usleep(1); // sleep the thread for a bit if there is no request
        }

        // get the next request
        pthread_mutex_lock(&mutex_request_queue);
        r = request_queue->next_request();
        DEBUG_PRINT("Service #: " << served << " (" << r->thread << " << r->track << ", " << r->time_offset << ")" << endl);
        // request_queue->queue_dump();
        pthread_mutex_unlock(&mutex_request_queue);

        // get the seek time
        seek_time = V + (r->track * M);

        // add the current seek time to the total service time
    }

```

ddr.photo

Mar 20, 06 16:52

ddr.photo

Page 3/8

```

        total_time += seek_time;

        // since we've successfully served one request, inc our counter
        served++;

        // sleep for the seek time
        SLEEP(seek_time);

        // set the request completed flag
        complete[r->thread] = true;
    }

    // set the done flag to signal the request processes to terminate
    done = true;

    // unblock all of the request threads that are still blocked
    memset(&complete,true,sizeof(bool) * N);

    // print out the average service time
    cout << "Average Service Time = " << (float)(total_time / served) << endl;
1;
//
    request_queue->queue_dump();

    // exit the thread
    return(0);
    pthread_exit(NULL);
}

void *request_thread(void *thread_id) {
    double delay = 0.0;
    double total_delay = 0.0;
    int total_requests = 0;
    request *r = NULL;

    while(!done) {
        // create a new request object
        r = NULL;
        r = new request;

        // setup the new request
        delay = (double)(RANDOM() * S);
        r->track = (int)(RANDOM() * T);
        r->time_offset = delay;
        r->thread = (int)thread_id;
        complete[(int)thread_id] = false;

        // add the delay time to the total delay time for this thread
        total_delay += delay;

        // sleep until the request needs to be posted
        SLEEP(delay);

        // insert the new request into the queue
        pthread_mutex_lock(&mutex_request_queue);
        request_queue->new_request(r);
        DEBUG_PRINT("Thread " << (int)thread_id << " created new request
(" << r->track << ", " << delay << ") " << endl);
//
        request_queue->queue_dump();
        pthread_mutex_unlock(&mutex_request_queue);

        // increment the total requests count for this thread
        total_requests++;

```

Monday March 20, 2006

Mar 20, 06 16:52

ddr.photo

Page

```

    }

    cout << "Thread id: " << (int)thread_id << "; Average Request Time:
(double)(total_delay / (double)total_requests) << "; Total Requests Made:
total_requests << endl;

    // exit the thread
    pthread_exit(NULL);
}
#ifdef __DDR_H__
#define __DDR_H__

/*****
* INCLUDES
*****/
#include <queue>
using std::queue;

#include <vector>
using std::vector;

#include <map>
using std::map;

#include <string>
using std::string;

#include <math.h>
#include <stdlib.h>
#include <unistd.h>

/*****
* MACROS
*****/
#define RANDOM() ((double)(rand() / (double)RAND_MAX))
#define DEBUG_PRINT(a) if(debug) { cout << a; }
#define SLEEP(a) usleep((useconds_t)(double)(a * (double)1000))

/*****
* CONSTANTS
*****/
const int N = 8; // number of processes
const float S = 120.0; // request time
const int T = 256; // number of tracks
const float M = 0.05; // track seek time
const float V = 3.0; // overhead time per request
const int REQUESTS = 10; // number of requests to run the sim for

/*****
* GLOBALS
*****/
bool debug = false; // set to true then
g info will print

/*****
* STRUCTURES
*****/

// a structure to hold information about a request to the disk drive
typedef struct request {
    int thread; // thread id of the request.

```

ddr.photo

Mar 20, 06 16:52

ddr.photo

Page 5/8

```

hread
    int            track;                // track # of the request
    double  time_offset; // time offset from the last request
};

// abstract base class for request queues
class RequestQueue {
public:
    // base constructor
    RequestQueue() {
        total = 0;
    }

    // base destructor
    virtual ~RequestQueue() {

    };

    // returns the next requests that should be processed
    virtual request* next_request() = 0;

    // returns the current # of requests pending
    virtual int request_count() = 0;

    // returns the total requests processed;
    int request_total() {
        return(total);
    }

    // adds a new request to the queue
    virtual bool new_request(request *r) = 0;

    // performs a dump of the current queue state
    virtual void queue_dump() = 0;

    virtual void print_stats() = 0;

protected:
    int total;
    string name;
};

// a request queue which implements a first-in, first-out priority
class RequestQueueFIFO : public RequestQueue {
public:
    RequestQueueFIFO() {
        name = "FIFO";
    }

    virtual ~RequestQueueFIFO() {
        // clean up any requests left in the queue
        while(!fifo.empty()) {
            delete fifo.front();
            fifo.pop();
        }
    }

    request* next_request() {
        request *r; // temp request pointer

        // get the top request and pop the value off the queue
        r = fifo.front();

```

Monday March 20, 2006

Mar 20, 06 16:52

ddr.photo

Page

```

        fifo.pop();

        // return the value of the pointer
        return(r);
    }

    int request_count() {
        // return the count
        return(fifo.size());
    }

    bool new_request(request* r) {
        // make sure we didn't get a null object
        if(r == NULL)
            return(false);

        // push the new request into the queue
        fifo.push(r);

        // increment the total counter
        total++;

        // well everything work so return something that represents
        return(true);
    }

    void queue_dump() {
        queue<request*> t = this->fifo;
        request* r = NULL;
        int i = 0;
        cout << "BEGIN QUEUE DUMP -----" << endl;
        while(!t.empty()) {
            r = t.front();
            t.pop();

            cout << "#" << i << ": request = (" << r->thread <<
            << r->track << ", " << r->time_offset << ")" << endl;
            i++;
        }
        cout << "END QUEUE DUMP -----" << endl;
    }

    void print_stats() {
        cout << "Queue Type: " << name << endl;
        cout << "Total Requests: " << total << endl;
    }

private:
    queue<request*> fifo; // STL queue class which is a FIFO queue
};

// A request queue which implements the SCAN priority algorithm
class RequestQueueSCAN : public RequestQueue {
public:
    RequestQueueSCAN() {
        head_position = 0;
        direction = 1;
        name = "SCAN";
    }

    virtual ~RequestQueueSCAN() {

```

ddr.photo

Mar 20, 06 16:52

ddr.photo

Page 7/8

```

        // clean up any requests left in the queue
        for(int i=0; i < T; i++)
            for(vector<request*>::iterator j=requests[i].begin(); j
!= requests[i].end(); ++j)
                delete *j;
    }

    request* next_request() {
        request *r;    // temp request pointer

        while(requests[head_position].size() <= 0) {
            head_position += direction;

            if(head_position < 0) {
                head_position = 0;
                direction *= -1;
            }

            if(head_position >= T) {
                head_position = T - 1;
                direction *= -1;
            }
        }

        // get the first element in the
        r = requests[head_position].front();
        requests[head_position].erase(requests[head_position].begin());

        // decrement the counter
        count--;

        // return the value of the pointer
        return(r);
    }

    int request_count() {
        // return the count
        return(count);
    }

    bool new_request(request* r) {
        // make sure we didn't get a null object
        if(r == NULL)
            return(false);

        // insert the request into the vector for that track
        requests[r->track].push_back(r);

        // increment the total counter
        total++;
        count++;

        DEBUG_PRINT(requests[r->track].size() << endl);

        // well everything work so return something that represents that
        return(true);
    }

    void queue_dump() {

```

Mar 20, 06 16:52

ddr.photo

Page

```

        void print_stats() {
            cout << "Queue Type: " << name << endl;
            cout << "Total Requests: " << total << endl;
        }

private:
    map<int, vector<request*> > requests;
    int head_position;
    int direction;
    int count;
};

#ifdef
FLAGS = -g
LIBS = -lpthread

ddr: ddr.o
    $(CXX) $(CXXFLAGS) $(FLAGS) -o ddr ddr.o $(LIBS)

ddr.o: ddr.cpp ddr.h
    $(CXX) $(CXXFLAGS) $(FLAGS) -c -o ddr.o ddr.cpp

.PHONY: clean
clean:
    @rm -f *.o ddr

1
Average Service Time = 9.915
Thread id: 2; Average Request Time: 51.7728; Total Requests Made: 4
Thread id: 0; Average Request Time: 74.3906; Total Requests Made: 3
Thread id: 6; Average Request Time: 42.8138; Total Requests Made: 5
Thread id: 3; Average Request Time: 34.9343; Total Requests Made: 6
Thread id: 5; Average Request Time: 58.642; Total Requests Made: 4
Thread id: 4; Average Request Time: 45.1473; Total Requests Made: 5
Thread id: 1; Average Request Time: 92.4354; Total Requests Made: 3
Thread id: 7; Average Request Time: 73.3606; Total Requests Made: 4
Queue Type: FIFO
Total Requests: 34

2
Average Service Time = 10.885
Thread id: 2; Average Request Time: 51.7728; Total Requests Made: 4
Thread id: 0; Average Request Time: 74.3906; Total Requests Made: 3
Thread id: 6; Average Request Time: 42.8138; Total Requests Made: 5
Thread id: 3; Average Request Time: 34.9343; Total Requests Made: 6
Thread id: 5; Average Request Time: 58.642; Total Requests Made: 4
Thread id: 4; Average Request Time: 45.1473; Total Requests Made: 5
Thread id: 1; Average Request Time: 92.4354; Total Requests Made: 3
Thread id: 7; Average Request Time: 73.3606; Total Requests Made: 4
Queue Type: SCAN
Total Requests: 34

```