

Synchronous Relaxation and Time Warp Algorithms: A Study

Mark Randles Edward Ondieki
randlem@bgsu.edu edwardo@bgsu.edu

Academic Supervisor
Hassan Rajaei, Ph.D

Department of Computer Science
Bowling Green State University, Ohio
07 Tuesday 2005

Abstract

There were a variety of ideas and goals pursued during the course of this project. First and foremost was an overriding goal of pulling the problem away from the system that it ran on. Previous implementations were dependent upon a certain architecture, and thus the system defined the problem size. This would need to change, the problem must become independent of the architecture. To do this, a variety of things were looked at. First we contemplated rewriting the initial work done on synchronous relaxation (SR). This accomplished it's goal of making the problem size independent of the architecture. However, there were still bottlenecks inherent in the system, so the use of Time Warp (TW) was studied. This also should accomplish the goal of making the problem size independent of the underlying architecture. Also discussed have been the use of logical processes and hybrid OpenMP/MPI mechanisms, on which work has not been done, but initial studies have been undertaken to study their feasibility.

Contents

1	Introduction	5
1.1	Monte Carlo Methods	5
1.2	Modeling Thin Film Growth Using Kinetic Monte Carlo	6
1.2.1	Kinetic Monte Carlo	6
1.2.2	Thin Film Growth	7
2	Synchronous Relaxation	11
2.1	Introduction	11
2.2	Original Algorithm	12
2.2.1	Boundary Events	12
2.2.2	Ghost Region	12
2.2.3	Data Structures	12
2.2.4	Functions	12
2.3	Improvements	14
2.3.1	Multiple Files	14
2.3.2	Readability	14
2.3.3	Function Overheads	15
2.3.4	Results	15
2.3.5	Analysis	15
2.3.6	Implementing the Program in C++	16
3	Time Warp	19
3.1	Introduction	19

3.1.1	Time Warp Algorithm: Classic	19
3.1.2	Time Warp Algorithm: KMC	20
3.2	Implementation	20
3.2.1	Lattice Class	22
3.2.2	MPIWrapper Class	23
3.2.3	Main Function and GVT	24
3.3	Issues	25
4	Future Work	27
4.1	Synchronous Relaxation	27
4.2	Time Warp	27
4.2.1	Debugging and Validation	27
4.2.2	Garbage Collection and Memory Management	27
4.2.3	Logical Processes	28
4.3	General	28
4.3.1	2D Decomposition	28
A	Synchronos Relaxiation Code	29
A.1	main.cpp	29
A.2	lattice.h	30
A.3	lattice.cpp	33
A.4	synch.cpp	50
A.5	comm.cpp	53
B	Time Warp Code	55
B.1	main.cpp	55
B.2	lattice.h	58
B.3	latconst.h	62
B.4	latprim.h	63
B.5	lattice.cpp	63
B.6	event.h	77
B.7	rewindlist.h	78

B.8	randgen.h	82
B.9	mpiwrapper.h	85
B.10	mpiwrapper.cpp	88
B.11	exception.h	94

Chapter 1

Introduction

In order to model physical processes like thin film growth on large areas and long time scales, more efficient parallel algorithms need to be developed. In this project we first studied a proposed conservative parallel algorithm known as Synchronous Relaxation(SR) algorithm to try and see whether it can be made more efficient. To do this, we rewrote the initial implementation of the algorithm porting from C to C++ . Though the program was slightly more efficient after revision it still was not scalable, that is, it lost its efficiency as more processors were added. We then proposed a new algorithm that would use a distributed simulation technique known as Time Warp. This optimistic algorithm should incur less synchronization costs than the SR algorithm and should therefore scale better

1.1 Monte Carlo Methods

Monte Carlo methods are numerical techniques often used to calculate integrals by using (pseudo) random numbers. A good example of this is the calculation of π using the following algorithm Monte Carlo for π

1. Generate i random number pairs (X_i, Y_i) where $0 < X_i < 1, 0 < Y_i < 1$.
2. Count number of pairs within the unit circle radius $r = 1$
3. $(count\ of\ pairs\ in\ r = 1) / (total\ pairs) \approx \pi/4$

[5]

Monte Carlo methods have also been successfully applied to solving various problems in statistical physics. A classic example is the Ising model used to study the

effect of temperature on a ferromagnet [3]. Initial Monte Carlo algorithms like the Metropolis algorithm worked by randomly selecting an event to occur and accepting or rejecting the event based on a criteria known as the Detailed Balance criteria. However such algorithms proved slow and inefficient for modeling systems at large sizes and time scale as is the case in thin film growth.

1.2 Modeling Thin Film Growth Using Kinetic Monte Carlo

1.2.1 Kinetic Monte Carlo

Kinetic Monte Carlo is a subset of Monte Carlo techniques that has proved successful and efficient in stochastic modeling non-equilibrium systems. Kinetic is derived from the n-fold way algorithm. In this method all the rates of possible events are known beforehand as well as the probability of each rate occurring. Events are then performed and system time incremented stochastically. Kinetic Monte Carlo is better suited for non-equilibrium systems like thin film growth or the Ising model because

1. It satisfies the Detailed Balance criteria an important property of stochastic processes.
2. It uses a finite amount of rates that is event rates are precalculated thus reducing computation
3. No events are rejected as opposed to previous algorithms like Metropolis algorithm which suffers from high rejection at low temperatures.

The algorithm follows the following basic steps

1. Generate a uniformly distributed random number $0 < X < 1$
2. Use the number X to select an event to occur depending on the probability of that event happening
3. Perform the event (e.g. flip atom or deposit monomer)
4. Generate a new random number $0 < Y < 1$
5. increment system time using the formula $T = T + \log(Y)/((eventrates))$
6. Repeat Step 1 until a stopping condition is equates true

1.2.2 Thin Film Growth

Thin film growth is the process by which a layer of atoms (monomers) are deposited on a surface (substrate) usually at low pressure and temperature so as to form a surface coverage that is a few atoms thick. The morphology of the thin film formed is highly dependent on various factors such as the rates of different events for example the number of atoms (adatoms, monomers) being deposited as opposed to the rate at which atoms on the surface (substrate) diffuse. Thin film growth also takes place in enormous time scales when considering events at atomic levels i.e atomic vibrations take place at nanoseconds yet it takes hours or minutes to grow a thin film device [2]. In our current model we studied a simple growth model known as the Fractal Growth Model. In this model atoms are deposited on the surface of an atom where they may then diffuse along the surface of the substrate. In case a diffusing (or deposited) atom encounters another atom, it reacts to form an island that binds the atoms to the surface. In case an atom encounters an island (two or more bound atoms) it will be captured by that island.

To model Thin film growth using KMC the basic data and data structures required are

1. A two dimensional lattice to act as a height map
2. A monomer list to keep track of the number of monomers as well as their positions on the lattice.
3. The surface diffusion rates as well as deposition rates

The algorithm proceeds like a standard KMC with the exception of steps 4 and 5.

1. Generate a random number $0 < X < 1$
2. Use the number X to select an event to occur depending on the probability of that event happening
3. Perform the event (e.g. deposit monomer)
4. Update the Lattice heights to record change
5. Update the Monomer list in case monomers were captured or added by event
6. Generate a new random number $0 < Y < 1$
7. Increment system time using the formula $T = T + \log(Y) / ((event\ rates))$
8. Repeat Step 1 until $(total_depositions / (dimension_x * dimension_y)) = 1$

```

Sample Output
g++ -c lattice.cpp main.cpp synch.cpp comm.cpp mpiwrapper.cpp
ndep=441
time=0.00556387
coverage=1.00227
time=0.00556387
Execution Time=2.2689
*****S*****
1 0 0 2 1 0 1 1 0 1 1 1 1 1 1 0 1 1 1 1 1 0
1 0 2 2 0 1 2 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1
1 0 0 2 2 2 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 2 2 1 1 1 1 0 0 1 1 2 2 2 1 2 2 1 1
2 0 1 1 2 2 2 1 1 1 2 2 1 2 1 0 1 0 1 2 2 1
0 0 2 2 2 0 1 1 1 2 1 2 2 1 1 3 1 1 2 2 0 0
0 0 1 1 2 2 1 1 0 1 1 0 2 2 3 1 0 1 1 2 0 1
1 1 1 1 1 0 1 1 1 1 1 1 1 2 1 3 1 0 1 1 0 1 2
2 1 1 2 1 1 2 1 1 1 1 0 1 1 1 2 1 0 0 1 2 1
2 0 2 2 1 2 2 2 1 1 1 1 2 0 1 2 3 1 0 1 2 1
1 0 1 1 1 1 2 2 1 0 2 2 1 0 1 1 2 2 1 2 1 1
2 0 1 1 1 1 1 2 0 1 2 2 2 1 1 1 1 1 1 1 2 0
0 0 1 1 1 1 1 2 1 1 2 0 1 1 1 1 1 1 1 1 0 0
1 0 1 1 2 2 2 1 1 1 1 2 2 2 1 0 1 1 1 0 1 0
2 0 0 2 1 2 1 1 2 1 0 2 0 1 0 1 1 1 2 2 2 0
2 0 1 0 0 2 2 0 2 2 2 2 1 0 0 2 2 1 2 2 2 1
1 1 2 0 2 1 2 1 2 2 1 1 1 1 2 1 1 2 0 1 1 2
1 2 2 2 1 1 2 2 2 2 1 0 1 1 2 1 2 1 1 0 2 3
1 0 2 1 2 2 1 2 2 2 2 2 2 2 2 2 0 2 1 3 1 1
4 0 1 1 1 0 1 2 2 1 1 0 1 2 2 2 2 2 1 1 4 6
mcount=1

```

Figure 1.1: Sample KMC Results

In Figure 1.1 the bold 3 indicates the presence of a sole monomer on the lattice. This is because it is higher than other heights around it and therefore it does not get bound. The monomer list contains the location (X, Y) pair of the monomer. In this case the monomer is of size 1 and the $(X, Y) = (9, 16)$.

Chapter 2

Synchronous Relaxation

2.1 Introduction

The first parallel KMC algorithm we studied was the synchronous relaxation algorithm (SR). The Synchronous Relaxation algorithm was first used to study circuit switched networks. Researchers like Lubachevsky then applied the algorithm to the Ising model and showed that are given a number of weak assumptions, the algorithm is efficient and scalable. The SR algorithm was later adapted by researchers at the University of Toledo to model thin film growth.

The Synchronous Relaxation algorithm is an optimistic algorithm which means that processors can process information independent of each other. The algorithm works as follows:

Step 1: Each Processor generates KMC events independent of other processors for a cycle time T

Step 2: When T expires each processor exchanges information on events that happened to see if any events generated by its neighbor affect it.

Step 3: If events in each processor do not affect events in neighbors process, all processors proceed to the next cycle

Step 4: If information received from one neighboring process affects another processor, all processors redo their KMC events incorporating information gathered.

Step 5: Processors again exchange information if events in neighboring processes affect other neighbors Step 4 is repeated. Or else all processors go to next cycle

[3]

2.2 Original Algorithm

2.2.1 Boundary Events

The main problem with parallelizing lattice based KMC problems like thin film growth is that of boundary events. Boundary events are those events that take place on the edge of the lattice and may affect the evolution of neighboring lattices. The SR algorithm is used to try and iteratively correct such errors.

2.2.2 Ghost Region

The ghost region is the area of a lattice that belongs to another lattice but whose events may affect the current lattice. In the initial implementation for the fractal model the Ghost region was one lattice step since diffusions take one step at time.

2.2.3 Data Structures

Modeling KMC events using Lattices can be rather complex. This is because a number of different data structures are required to keep track of the system.

h[x][y]: 2x2 Integer array for the height map. (Appendix A.2)

indexa[size], ipointa[size], indexc[size]: Three lists to keep track of mobile monomers on the lattice. (Appendix A.2)

myeventlist[size*size]: An event list to track. (Appendix A.2)

bdyeventlist[size]: Keeps track of boundary events generated. (Appendix A.2)

ranlist[size]: list of random numbers generated. (Appendix A.2)

2.2.4 Functions

DoKMC: using a counter this function selected the next random number from the list and used this number to determine which event should take place (diffuse or deposit)

CalcTime: calculates the time increment $DT = \log(Rand)/\text{sum}(\text{eventrates})$ and increments the time. It is always called after DoKMC().

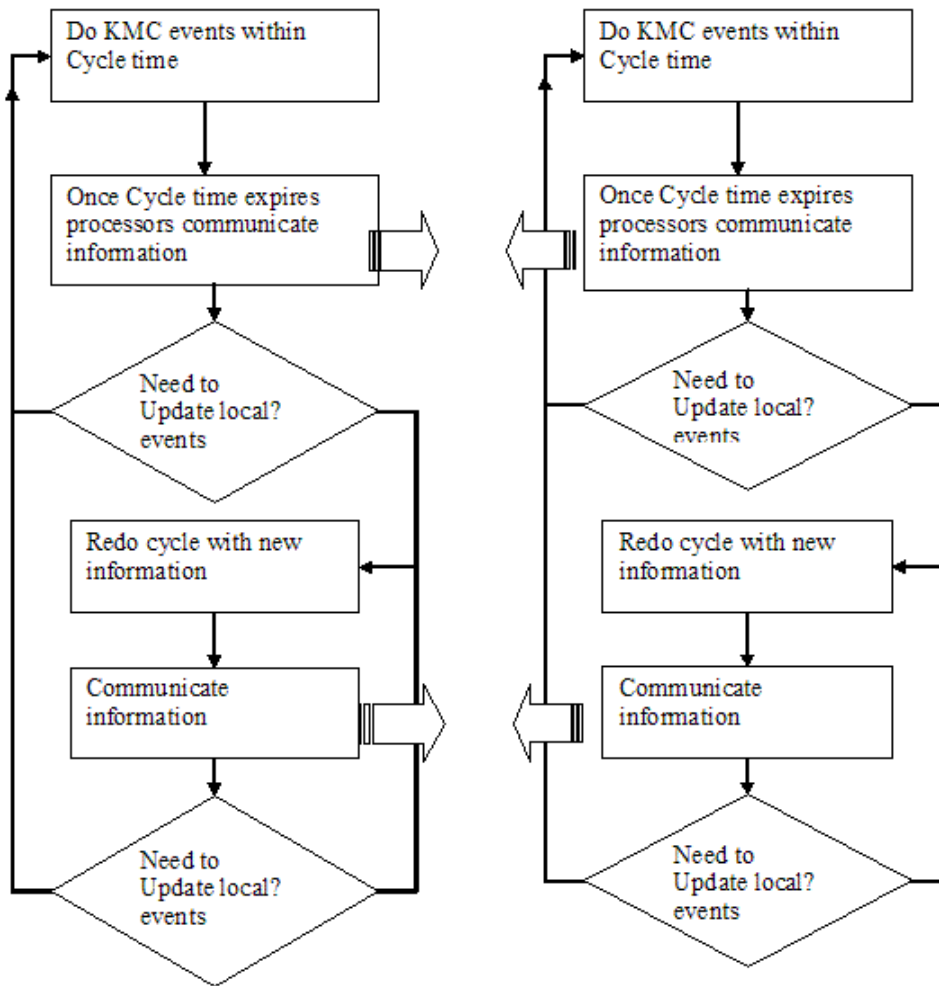


Figure 2.1: Synchronous Relaxation Flow Chart

Diffuse: This function was called by the `DoKMC()` function when a diffuse event was selected. The function selects a random location on the lattice and increments it by one. It then calls `upnbhd()`.

Deposit: This function was called by the `DoKMC()` function when a deposit event was selected. The function selects a random location on the lattice and increments it by one. It then calls `upnbhd()`.

Upnbhd: When a monomer moves (diffuses) or is first deposited this function is called to check whether there are any neighboring monomers or islands that would bind the monomer to the lattice. If there are no monomers or islands surrounding the current monomer then it is added to the list. If there is a monomer/island next to the list then the monomer is deleted from the list or not added in the case of a deposition. If the monomer binds with another monomer that other monomer too will be deleted.

Undoevent: This is used in the SR algorithm to undo the events that occurred and restore the list and lattice their state at the beginning of the cycle.

BufferSendrecv: This function communicates boundary events to neighboring processes.

UpdateBuffer: This function uses boundary events received to correct the lattice. It takes the boundary events received at executes them in order of the time.

2.3 Improvements

2.3.1 Multiple Files

The initial algorithm was implemented in a single C file with a very large number of global variables as well as a large number of functions. To abstract different aspects of the program such as communication functions like `buffersendrecv` or lattice manipulation functions like `diffuse` or `deposit`, we split the program into different files. The function headers were contained in the header files such as `lattice.h` (Appendix A.2) and implemented in `lattice.cpp` (Appendix A.3) and `comm.cpp` (Appendix A.5 files).

2.3.2 Readability

In order to work on the existing code we had received we needed to get it more readable. To make it more readable, we first eliminated all global variables. We did this by encapsulating them in a `Lattice` structure and then passing this structure to functions by reference. This allowed us to follow the execution of the program easier as well add new capabilities to it.

Execution Time with function overhead (Sec)	Execution Time without function overhead (Sec)
138.16	109.12

Table 2.1: Execution Time vs. Function Overhead

Execution Time without rewind list (Sec)	Execution Time with function overhead (Sec)
138.16	119.12

Table 2.2: Execution Time vs. Rewind List usage

2.3.3 Function Overheads

One of the issues slowing down the execution of the program was excessive function calls. This was taking place especially in the `Upnbhd()` function where a single call from `Upnbhd()` would call a helper function eight or more times. By making the small functions inline this helped reduce execution time. Rewinding to first boundary event

One of the main problems with the initial implementation is that when the program entered the correction phase it had to return to the beginning of the cycle as opposed to correcting after the first boundary event. This therefore increased the computation time and cut the parallel efficiency of the program. In order to make it possible for the program to rewind to the first boundary event, all changes made to the monomer lists would have to be recorded and then undone. This meant adding extra data structure known as `ListChange[]`. `ListChange[]` kept track of all the insertions and deletions that took place in the monomer list. `ListChange[]` also kept track of time each of these changes occurred. When the program synchronized and tried to correct errors from boundary events, it would loop through the `ListChange[]` structure and restore the list to the state it was at the time of the first boundary event.

2.3.4 Results

The results from the improvements we made were mixed. When function overhead was reduced, the execution time was reduced. The results are contained in Table 2.1

The rewind to first boundary events did not increase parallel efficiency or reduce execution time due to the complexity of manipulating the three monomer lists. This complexity made the procedure complex and error prone. Execution time was actually increased due to these changes. The results are contained in Table 2.2

2.3.5 Analysis

Overall the Changes made to the initial program did not adequately improve the performance of the implementation. The main factors that controlled efficiency where the cycle

Execution Time in Sec Cycle Time=10e3 Diffusion/Deposition=10e5	Execution Time in Sec Cycle Time=10e6 Diffusion/Deposition=10e5
138.16	90.12

Table 2.3: Execution Time and Cycle time

time, lattice size and diffusion/deposition rates. For example reducing the cycle time to 10e-6 s with a diffusion/deposition rate of 10e6. The results are contained in Table 2.3

Parallel efficiency was difficult to measure in the initial program because we relied on the OSC cluster and any requests for more than 4 processors resulted in a long queue that could take hours or weeks before a turnaround.

2.3.6 Implementing the Program in C++

Motivation

The motivation for moving the algorithm from C to C++ where

1. It allowed for a much better abstraction of the program
2. Lattice object would allow for process-processor independence
3. The Lattice class could be reused easily in other algorithms like Time Warp
4. Easier to debug

Lattice Object

The lattice Object was now made up of an array of Sites. Each site contained information about its location, height and if a monomer was present the location of that monomer on the monomer list.

Code: Site and Lattice Object

```

Class Site
{
    int index;//location of monomer on the monomer list
    int height;
    point position;
}

Site Lattice[X][Y];//lattice object.
```

Num Processes	Execution Time	Parallel Efficiency
1	11.1	1
2	9.233	0.6
4	8.7	0.3
5	8.66	0.25
8	10.38	0.13
10	10.23	0.10

Table 2.4: Execution Time and Parallel Efficiency

The main advantages of this approach was that it simplified the updating of monomer lists. This meant that we could implement a rewind list easier since there is only ONE list to worry about.

Multiple files

Moving from C to C++ allowed us to further abstract the code. The lattice functions could now be separate from the communication functions. All the MPI calls were encapsulated in an easy to use `MPIWrapper Class`.

Rewind List

As mentioned earlier the rewind list class was now easier to implement since we had a single monomer list.

Results and Analysis

Using a 800 by 200 lattice we tested the parallel efficiency of the algorithm on the local cluster the results are contained in Table 2.4. They clearly show that the performance degrades at nearly exponential rate as the number of processes increases. This is due to the conservative nature of the algorithm, each synchronization cycle taking longer as more processes are added. Ultimately there is a hard performance limit, as eventually the communication time will significantly outweigh any possible benefit from parallelization.

It may be possible to tune the SR algorithm with some sort of emergent heuristic algorithm to reduce or extend the cycle lengths between conservative synchronizations. However, this would add further complication the the algorithm for what could be seen as a negligible benefit, as you're still being restrained by the conservative nature of SR.

Chapter 3

Time Warp

3.1 Introduction

During the course of this project there were two focus of studies on Time Warp (TW). First off we must define what it means to be a TW system and secondly we must configure the KMC problem in such a way that it fits in with the TW algorithm. The first study was rather simple in it's function. However the second was much more in-depth and was the primary focus of work.

3.1.1 Time Warp Algorithm: Classic

Classically, TW is defined as a "*optimistic*, object-oriented, distributed, discrete-event simulation mechanism" [1]. This is different then classical approaches to distributed-event simulation (DES), in that it's not conservative. This allows TW to better exploit the "concurrency in distributed simulations" [1].

The distinction must be made between optimistic and conservative synchronization algorithms. Time Warp falls into the former category, which is to execute and attempt to solve for the solution correcting errors as they occur, which is an asynchronous operation. In a conservative algorithm, part of a solution is solved, after which any errors are corrected, ad infinitum until the correct solution for that subsection of the entire solution is correct. Then the algorithm can move on to the next subsection. The entire algorithm is synchronous, as all executing programs will proceed in lock-step until they arrive at the final solution.

In a TW system you have three components. The first component is the local virtual time (LVT), which is local to each TW object, and is the method by which that object timestamps and synchronizes messages. The LVT advances independently of each other clock, this is the asynchronous nature of the algorithm. The system is optimistic because each TW object advances even if there is the possibility for an error in the computation.

The second component is the error correction or *rollback*. This move the simulation back in time to a point before the error occurs. The simulation can then proceed from this point, fixing the error when it occurs. Any computations (events) that are not performed when the error is corrected are canceled by anti-messages.

However, there is a trade-off for the optimistic nature of the algorithm. We incur a penalty each and every time we must roll back the simulation to fix an error. The penalty is in the cost of the the roll-back and the extempores calculations that we may have performed.

Lastly, there is the third component, the global virtual time (GVT). This is the lowest LVT that is associated with any TW object. The GVT is used to garbage collect out-of-date objects, messages, and roll-back checkpoints.

3.1.2 Time Warp Algorithm: KMC

Applying TW to the KMC simulation that was the subject of this project was rather straightforward. Already existing in the problem are the concepts of LVT and error correction. The concept of GVT was implicit in the SR error correction, since the SR would reuse buffers during each cycle, with is approximate with the effects of GVT [5].

The only concept that does not carry directly over from traditional TW is the concept of an event. In DES, events are scheduled for the future, based on the current event. In KMC, we do not know any future events, and only know the current event. Therefore, a KMC event is the computation of the current event, and not the scheduling of a future event. This doesn't effect the idea of the asynchronous clocks of TW nor the idea of error correction, but it is important to understanding the caveats of implementation and to sort out the meaning of the terminology.

The mechanics of the KMC algorithm have already been described in Section 1.2.2. Since the goal of this project was to delve into the underlying mechanics of the algorithm, of which the KMC was an appropriate problem, the code needed to be written in such a way that it was both flexible and robust. To facilitate this it was determined that the C++ language would be up to the job. It was also determined that any of the KMC parts should be as self-contained as possible, which like the re-implementation of SR, was in the form of a Lattice class. The Lattice class was supported by various layers 3.1 of abstraction that removed as much of the implementation dependence as possible.

3.2 Implementation

The actual code base for this implementation of the TW: KMC algorithm is split into a number of files, both C++ source and header files, which contain a number of classes that do nearly all the work. A text dump of the source files are available in Appendix B.

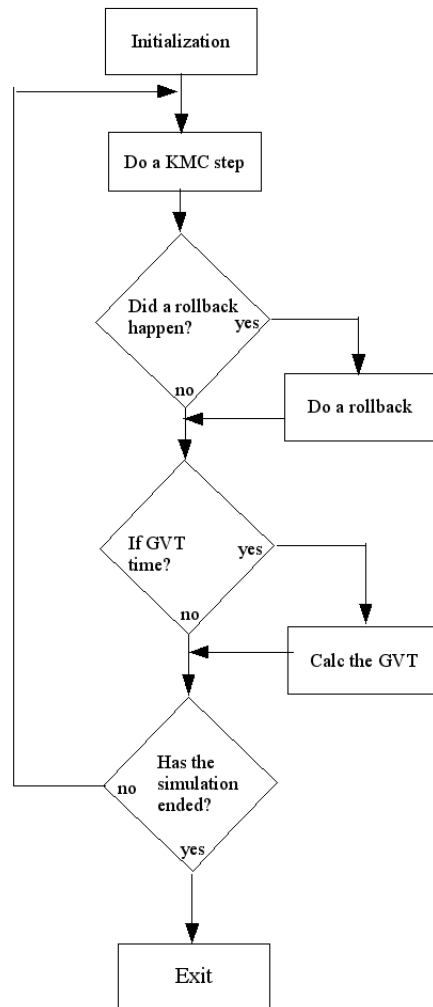


Figure 3.1: Sample KMC Results

3.2.1 Lattice Class

This sub-section refers to the following appendix sections: B.2, B.5, B.3, B.4, B.7, B.8, B.6.

The core of the TW: KMC implementation is a Lattice class. The design of this class was twofold: 1) To improve on the previous Lattice class as implemented in the new SR code (Appendix A) and 2) To build in a new way of handling events that improves the storage and computational efficiency of error correction.

Before the Lattice class proper can be addressed, one must understand the breakdown of the various parts. In the file `latprim.h` (Appendix B.4) are the primitive data structures that make up the most atomic parts of an event. In the file `event.h` (Appendix B.6) is the definition of the Event class. The file `latconst.h` (Appendix B.3) contains all of the various lattice constants, such as deposition rate and lattice size. The two support classes, a RewindList abstract data structure (Appendix B.7) and a rewindable random number generator (Appendix B.8), are contained in separate files. The Lattice class is contained in the files `lattice.h` and `lattice.cpp` (Appendix B.2 and Appendix B.5, respectively).

The core of the Lattice class is the event processing. In previous algorithms, as events are generated they're processed immediately. Remote events are stored in a buffer and exchanged during a sync. The same holds true in this implementation, but it is done in a more object-orientated way, where the event becomes an object, and the object is worked upon.

The first that that must be done is determine the locality of the event that must be processed. Since we're going to be doing error correction using remote events, we must make sure to process events from the remote queue in the correct order. To do this, we calculate the time of the next local event and compare that value with the earliest event in the remote event queue. The lowest event is processed. If it is a local event, we create the event. For remote events we grab the event out of the queue.

Once we have the event, we proceed to the next step. This step is to commit the event to the lattice. At this stage we don't care if it's a remote event or not, and we perform the action, wither deposition or diffusion, After the event is committed we push it onto a stack of events processed.

Remote events are passed to the communication class, `MPIWrapper`, during this step as well. This is a pretty abstract process from the Lattices point-of-view, as all the Lattice needs to know is which boundary the Event occurs on, and what the event is. The wrapper layer (in this case `MPIWrapper`) can then translate the passed data and send the event. The details of this operation in the `MPIWrapper` layer is covered in Sub-Section 3.2.2.

The Lattice class also needs to check for new remote events. Currently this is done periodically in the main function, after one event is processed. This is important, since we must gather the remote events to correct any error we may have made, and to have the events on hand so we don't make an error in the future. Some though should be put into how often we poll the queue for new remote events. More often then not, polling once a cycle, we're

not going to have any remote events. However, if we make the polling interval too long, the chances for errors to occur goes up drastically, which means more rollbacks.

The function that polls for remote events is responsible for detecting errors in the simulation as well. When it finds an error, it will trigger a rollback. The rollback function takes a single parameter which is the time at which the error occurs. The rollback function loops through the event stack and compares the top of the stack with the passed timestamp, stopping when the top of the event stack is less-than the passed timestamp. Each cycle of the loop, if we're not going to halt, the top event is popped off and the event is rolled back on the lattice. If the event is a remote event, it's stored in the remote event queue, otherwise the event will be discarded. Sadly, this means that we must send anti-messages for any event that was sent as a boundary event. Normally this would not be an issue, as the TW algorithm would use lazy evaluation in sending anti-messages, only sending anti-messages for those events that would never happen again. However, this will be a rare, if not impossible situation, as any KMC event influences the rates for future events, and therefore the probability of that event happening in the same way again. Since we are almost guaranteed to not repeat the event, we might as well cancel all of the sent events with their anti-message, so we don't need to incur the overhead of lazy evaluation.

The Lattice class would be nearly useless without data output. There are a variety of ways that one can get data from the Lattice class, including dumps of the height map in both text and as a PNG file. Also the Lattice can dump its current monomer list and a collection of various counts and statistics that describe the simulation to that point. Most of these functions take an output stream as a parameter, with one taking the file name that should be created.

3.2.2 MPIWrapper Class

This sub-section refers to the following appendix sections: B.9, B.10.

The MPIWrapper class is divided into two files, the definition file `mpiwrapper.h` (Appendix B.9) and the implementation file `mpiwrapper.cpp` (Appendix B.10). Included in these files are any helper structures, such as `message`. As the name implies, this class is a wrapper for all of the MPI related calls the program must make.

MPIWrapper provides a number of interfaces for doing various helpful things with MPI. Included are calls to do global scatter-gather operations and do barrier calls. The most important interface provided is the sending of events between processes.

The event passing interface is very abstract, compared with previous implementations. In the original SR code, the MPI calls were performed directly, so the lattice functions needed to know the particulars of the MPI calls and such. With the MPIWrapper, the Lattice class only needs to pass the event as an Event class pointer and the direction (or boundary) direction as an enumerated integer. The MPIWrapper call transforms the object as passed into a data structure (type `message`) which it then can send to the corresponding node. To mark the type of message, either an event message or anti-message, the tag property of the send

call is used.

It should be noted that the send call used in this operation is `MPI_Bsend()`, which buffers the message in a local buffer (for this project a 10MB buffer), and only sends the message when the receiving node is ready to receive. This allows for a asynchronous send, with a blocking receive.

Since the Lattice class must also gather remote events, there is also an interface that returns a list of events as they were received. There is the possibility that the wrapper could stall out program execution if there are many, many events waiting to be received. However, this would require a enormous number of events, and in such a situation the current communication model would not be ideal.

The interfaces for sending and receiving anti-messages are the same as the ones provided for normal messages. The only difference between the function is the tag passed to the MPI calls within.

Also exposed are a suite of function that perform a `MPI_Allreduce()`, each function uses a different data type, and takes a data value and a MPI operator as parameters. This is useful for collecting global statistics and performing the GVT calculation.

Lastly, the MPIWrapper exposes a method to output statistics about it's communications, such as counts of events sent and received and pertinent data about the cluster itself such as number of nodes and who are the neighboring nodes.

3.2.3 Main Function and GVT

This sub-section refers to the following appendix sections: B.1.

The main function of the project is contained in the file `main.cpp` (Appendix B.1). Laid out in the main function is the highest level of abstraction in the code, along with the GVT calculation and the calculation of the stopping condition, which is currently the coverage of the lattice.

The level of abstraction in `main()` is meant to insulate the top level code from the actual processes underneath. The program is essentially constructed of a main loop that calls a function to process the next event, and then receive and processes any remote events. Also, at periodic intervals it performs a GVT calculation. Currently this does nothing as the garbage collection has not been implemented. The `main()` function is responsible for outputting the data in various forms along with logging the progress of the simulation for debugging purposes.

The GVT calculation is an all reduce operation that finds the minimum LVT of all the TW objects in the simulation. This value is set as the GVT. No event will ever come before this time, so we can effectively delete any data up to that point. However, this is not done right now. GVT also makes a decent exit condition for the simulation, as we need not compare the results against previous implementations.

To compare the simulation results against other implementations we need to find the global

coverage of the simulation, which requires a second all reduce operation, this time to sum the number of depositions and therefore the number of monomers in the simulation. To find the coverage we divide the number of monomers by the total surface area of the lattice. We exit when that value grows greater than 1.0 (100%).

3.3 Issues

A variety of problems were encountered while implementing the TW algorithm previously described. Most have been solved, many being bugs with the underlying communication system doing things that more than likely were not intended by the authors. However, at the time of writing, there has been a bug that shakes the very root of the problem. Work has stalled out until this bug can be over come.

The bug in question deals with the validity of results obtained at the completion of the simulation. The results are not consistent from one simulation to the next, which indicates a problem with the rigorous nature of the algorithm. Normally, the simulation should always, *always*, produce the same results for a given input, barring major hardware failure or a bug (such as the Pentium Floating Point bug) in the actual system hardware. However, this bug is not a hardware bug, but rather a software bug of the systemic nature.

There could be a variety of causes for such a bug, including a logic error in the error-correction mechanism (rollback) to a race condition. So far, only a few ideas have been floated, mostly dealing with a race condition between processes to get the first event. This race condition would then influence the results of the simulation enough to produce drastic results to the final output, as is the nature of the KMC algorithm, one event precipitate further events of random nature. Logically, however, there are only a few conditions that such a race condition could cause, as there is a limited number of "first events", to start the ensuing chain leading up to the final result. Suppose that there are n processes, each operating independently of each other. In the case that thread #1 were to generate the first event, we would see one result, if the rest of the simulation were truly rigorous. The same for thread #2 generating the first event. It follows that for n processes, we get n results. This is not observed.

Even if there is not a race condition currently, it may be adding to the problem of debugging the actual cause of the differing results. A sort of error on-top of an underlying error. One thing that this implementation discards is the idea of a circular lattice, which is present in the original and rewritten SR algorithms. The author did not think of the implications of this decision, namely events that happen on the two boundaries that have no neighbors. In the current TW implementation these messages are just discarded in the wrapper layer. However, they should be "wrapped around". This would cause the simulation to disregard the race condition, as it does not matter which process sends or receives the first message since all messages are handled. The only difference between runs should be the order of the lattices, where the results will be transposed over some number between 0 and $n - 1$ offset from a null state, as the lattice is now a ring, rather than a bar.

It would be more probable that the error occurs during the error-correction sequence. A logic error in this section of code could precipitate a chain of errors which would end up drastically changing the results. However, this is a difficult thing to debug, not only because of limitations in the work environment, but because there are many, many events exchanged for a given time period, and it is difficult to sift through them all and make sure all of the events were handled. This is the next task to tackle and should take a good long time to complete. Not only will further code need to be written to facilitate this debugging, but then the code will need to be corrected, which in the extreme possibility could necessitate a total rewrite of the error-correction code. However, this is in the extreme case that it's totally wrong, which isn't very likely. An educated guess would be that the conditions on the while loops that control the correction are wrong.

It should be noted that both of these conjectures about the possible cause of the current major bug are just that, pure speculation. Lots of time has been spent on this algorithm, so they are educated guesses, but the root cause could turn out to be something entirely different from the ideas laid out here. However, once this bug is resolved, there is a good chance that the algorithm is correct and hard performance results can be gathered and future work can take place.

Chapter 4

Future Work

4.1 Synchronous Relaxation

We feel that there is little work left to do on the SR algorithm. The promise of our work is in the potential improvements that could be made with the Time Warp algorithm. However, should Time Warp fail, this would be the most promising path for further research, so any work should not be scrapped, but rather shelved until a later date.

4.2 Time Warp

4.2.1 Debugging and Validation

This is by far the first and most important thing that will need to be done to this algorithm. Only once this step is completed can any future work proceed. To do this first the difference in results between runs must be identified and suppressed or eliminated. Secondly, the validity of the results should be sought. This step could prove to be difficult, as the original SR algorithm should always produce different results than the TW algorithm. Because of this, the validity of results will nearly always been in question, and should be studied with veracity until we are certain. Time to complete: 10-30 hours.

4.2.2 Garbage Collection and Memory Management

Right now, there is no garbage collection implemented in the TW algorithm. Since this algorithm is very memory intensive, garbage collection would be a good thing. To do this we would need to call a routine every time the GVT is calculated that will delete all of the events that occur before the GVT. This will free memory for the next GVT cycle. Also, if user-space memory management was implemented, we could reuse the memory to keep

from reallocating the same chunks thus providing more speed to the algorithm. Time to complete: 10-20 hours.

4.2.3 Logical Processes

One of the primary goals of the TW algorithm was to add more processor independence to the KMC algorithm. Thus far, the algorithm has seemed to do this, and now the limiting factor is the problem size and physical hardware restrictions such as memory. However, we can go a step farther and apply a concept called logical processes to the algorithm to produce a truly processor independent implementation.

Logical processes essentially encapsulate a single KMC process. Currently, each node on a cluster will execute one process. With logical processes each node will execute tens, hundreds, or thousands of similar processes in such a way as to simulate the running of just one process on a single node. To accomplish this, there must be a middle-ware layer between the program and the message passing interface. This middle layer handles both external MPI communication with other nodes, and internal node communication between logical processes on the same node. The goal of this middle layer is to make it as transparent as possible, the calling thread should not know if the process that will be receiving the message is local logical process or a remote (on another node) logical thread.

To complete this task a variety of things will need to be completed. First a formal study of the middle-ware layer as exposed by the BGTW simulation library should be undertaken to judge the usefulness of this established code base. Secondly, any changes that will need to be made to the algorithm should be performed and then the simulation revalidated to ensure that the algorithm is still correct. Thirdly, the middle-ware layer should be integrated into the existing algorithm and validated. This last step will produce the final logical process version of the TW algorithm. Time to complete: 20-50 hours.

4.3 General

4.3.1 2D Decomposition

The current algorithms work with at 1D decomposition or strip lattice where there is at most only two neighboring lattices any given lattice. This is limiting however in total scale of problem, as to increase the total area we must increase the height and width of the lattice. Eventually the number of boundary events exchanged over a single boundary will slow the overall execution down. At this point it will be beneficial to split a single sub-lattice into a 2D map of adjoining lattices. We can then scale the problem up further till we reach the same limiting factor as before. Time to complete: 20-40 hours.

Appendix A

Synchronos Relaxiation Code

A.1 main.cpp

Code: main.cpp

```
#include "lattice.h"

MPIWrapper mpi;
int main(int argc, char * argv[])
{
    int i, ctr, myid=0;
    int totaldep;

    mpi.init(&argc, &argv);

    myid=mpi.getRank();
    Lattice newlatt;
    float cov, COV=1, T=0.001;
    newlatt.nbhr[left]=LEFT(myid);
    newlatt.nbhr[right]=RIGHT(myid);

    while(cov<=COV)
    {
        /**/
        newlatt.time=0;
        newlatt.nevent=0;
        newlatt.randgen();
        newlatt.iran=0;
        newlatt.saveconfig();

        newlatt.bdycountrec[left]=0;
```

```
newlatt.bdycountrec[right]=0;

while(newlatt.time<T )
{
    if(newlatt.time<=T)
    {
        newlatt.doKMC();
        newlatt.calctime();
    }
    newlatt.savebdylist();
    sendmsgs(&newlatt);
    synch(&newlatt);
    mpi.allReduce(&totaldep,&newlatt.ndep,1,MPI_FLOAT,MPI_SUM);
    cov=(float) totaldep/(float) (2*size*size);
    cout<<"cov="<<cov<<endl;
    ctr++;
}

cout<<"latone*****"<<endl;
newlatt.p();

mpi.shutdown();

return 0;
}
```

A.2 lattice.h

Code: lattice.h

```
//class node;
#ifndef LATTICE_H
#define LATTICE_H
#include<iostream>
using std::cout;
using std::endl;
```

```

#include <math.h>
#include <stdlib.h>
#include "mpiwrapper.h"
#include "boundaryevent.h"

const int size=20;
const int latsize=size*size;
const int dir=8;
const int left=0;
const int right=1;
const int update=0;
const int diffevent=1;
const int depevent=2;
const int Np=2;
#define LEFT(a) (((a - 1) >= 0) ? (a - 1) : (Np - 1))
#define RIGHT(a) (((a + 1) < Np) ? (a + 1) : 0)

class kmcevent{
public:
    site oldsite;
    site newsite;
    int ranseq;
    double t;
    int tag;
};

class slist{
public:
    point monomerloc[size*size];
    int mcount;
    int ndep;
};

struct listofchanges{
    int oldsite;
    point oldval;
    int newsite;
    point newval;
    int tag;
    float time;
};

class Lattice
{
private:
    //lattice array
    site location[size+2][size];
    //monomer list

```

```

    point monomerloc[size*size];
    //diffusion and deposition rate
    float deprate,difrate;
    bool event;
    //monomer count and number of depositions and random number count
    int mcount;
    //directions to move
    point mdir[dir];
    //random move function
    point ranmove(site);
    //check if site is bound
    int getbonds(site,point * bondpt);
    //return neighbors of site
    int getnbhrs(site,point * bondpt);
    //delete monomer from list
    void deletemonomer(point mysite);
    //keep track of changes to list
    void addmonomerchange(int tag);
    //check if deposited/diffused monomer is next to a monomer
    bool neighborIsMonomer(point pt,site mysite);
    listofchanges change[1000];
    void restorelist(float t);
public:
    Lattice();
    ~ Lattice();

    kmcevent myeventlist[10000];
    slist oldlist;
    //check neighborhood of monomer
    bool checkupdatebonds(site mysite);
    void deposit();
    void diffuse();
    void doKMC();
    //print

    /*Next level*/
    void addbdyevent(site,site,float,int);
    /*
        void sendrecv();
        void updateLattice();
        void restoreLattice();
    */
    void restoreLattice();
    void undoevent();
    void saveconfig();
    void restorelist();
    void restorelist(float time);
    void savebdylist();
    int comparebdylist();

```

```

void p();
void calctime();
void upnbhd(site);
void checksite(site);
void updateBuffer(int iranflag);
void sorting_nbevent();
float time,T;
void randgen();
float ranlist[10000];
boundaryevent bdyevent[2][500],bdyeventrec[2][500],olddbdyeventrec[2][500],
               sortbdyevent[500];
int bdycount[2],olddbdycountrec[2],bdycountrec[2];
    int myid,nbhr[2],changecount;
int redoflag,nevent,nupdate,iran,ndep,subcycle,nK,tnbdyevent;

};

void synch(Lattice * newlatt);
void sendmsgs(Lattice * newlatt);

#endif

```

A.3 lattice.cpp

Code: lattice.cpp

```

#include "lattice.h"
point Lattice::ranmove(site mysite)
{
    point pt;
    pt.x=0;pt.y=0;

    ////cout<<xdir<<" ";
    int prob=ranlist[iran]*4;//rand()%4;
    iran++;
    pt.x=mysite.pos.x+mdir[prob].x;
    pt.y=mysite.pos.y+mdir[prob].y;
}

```

```

//Allow boundary movement
//No!!!
    if(pt.x>=size+2) pt.x-=1;
    if(pt.x<0)pt.x+=1;
        if(pt.y>=size)pt.y-=2;
    if(pt.y<0)pt.y+=2;
    return pt;
}

Lattice::Lattice()
{

    float ratio=0.0;
    float prob;
    point newsite;
    mcount=0;
    deprate=1.0f,difrate=1.0e3;
    ndep=0;
    nevent=0;
    time=0;
    iran=0;
    subcycle=0;
    T=0.001;

    int i=0,j=0;
        for (i=0;i<2;i++)
            bdycount[i]=0;

    for (i=0;i<2;i++)
        bdycountrec[i]=0;

    for (i=0;i<2;i++)
        oldbdycountrec[i]=0;

    for (i=0;i<size+2;i++)
    {
        for (j=0;j<size;j++)
        {
            newsite.x=i;
            newsite.y=j;
            location[i][j].pos=newsite;
            location[i][j].h=0;
            location[i][j].index=-1;
        }
    }

    //generate random numbers
    //randgen( );
    /*

```

```

    generate points in the form e.g (1,-1) move one unit left in x
    and and one unit up
    do not allow particle to remain (0,0)
    */

    /*Set directions*/

    mdir[0].y= 1;  mdir[0].x= 0;
    mdir[1].y= 0;  mdir[1].x= 1;
    mdir[2].y=-1;  mdir[2].x= 0;
    mdir[3].y= 0;  mdir[3].x=-1;
    mdir[4].y= 1;  mdir[4].x= 1;
    mdir[5].y= 1;  mdir[5].x=-1;
    mdir[6].y=-1;  mdir[6].x= 1;
    mdir[7].y=-1;  mdir[7].x=-1;

}
void Lattice::calctime()
{
    float    monomer=(float)mcount,
            Drate=difrate*monomer*0.25f,
            N=(float)latsize,dt,prob,
            totaldep=deprate*N;
            prob = ranlist[iran];
            iran++;
            dt=-log(prob)/(Drate+totaldep);
            time+=dt;
}
Lattice::~Lattice()
{

}

int Lattice::getbonds(site mysite,point * bondpt)
{
    int ctr=0,i=0;
    point pt;
    pt=mysite.pos;
    for(;i<dir;i++)
    {
        pt.x+=mdir[i].x;
        pt.y+=mdir[i].y;
        if((pt.x<0)
            ||(pt.x>=size+2)
            ||(pt.y<0)
            ||(pt.y>=size))

```

```

    {
    }
    else
    {
        if(location[pt.x][pt.y].h>=mysite.h)
        {
            bondpt[ctr]=pt;
            ctr++;
        }
    }
    pt=mysite.pos;
}
return ctr;
}

int Lattice::getnbhrs(site mysite,point * bondpt)
{
    int ctr=0,i=0;
    point pt;
    pt=mysite.pos;
    for(;i<dir;i++)
    {
        pt.x+=mdir[i].x;
        pt.y+=mdir[i].y;
        if((pt.x<0)
            ||(pt.x>=size+2)
            ||(pt.y<0)
            ||(pt.y>=size))
        {
        }
        else
        {
            bondpt[ctr]=pt;
            ctr++;
        }
        pt=mysite.pos;
    }
    return ctr;
}

void Lattice::deletemonomer(point pos)
{
    point lastpt;
    if(mcount>1)
    {
        lastpt=monomerloc[mcount-1];
        addmonomerchange(diff,lastpt);
        location[lastpt.x][lastpt.y].index=location[pos.x][pos.y].index;
        monomerloc[location[pos.x][pos.y].index]=monomerloc[mcount-1];
    }
}

```

```

        location[pos.x][pos.y].index=-1;
        mcount--;
    }
    else
    {
        location[pos.x][pos.y].index=-1;
        mcount--;
    }
}

bool Lattice::neighborIsMonomer(point pt,site mysite)
{
    bool bMonomer=false;
    //if index is not -1 and same height as my location
    if (location[pt.x][pt.y].index!=-1 &&
(location[pt.x][pt.y].h==location[mysite.pos.x][mysite.pos.y].h))
    {
        bMonomer=true;    }
    return bMonomer;
}

bool Lattice::checkupdatebonds(site mysite)
{
    /*
    Possible Scenarios
    Deposit
    1. Monomer encounters no cluster or other neighbor monomer
    -No bonds
    2. Monomer encounters cluster
    -bond delete monomer from list
    3. Monomer encounters single monomer
    -bond. delete BOTH from list

    Diffusion
    1. Monomer encounters no cluster or other neighbor monomer
    -No bonds
    2. Monomer encounters cluster
    -bond delete monomer from list
    3. Monomer encounters single monomer
    -bond. delete BOTH from list
    */
    bool bond=false;
    point pt, bondpt[dir],lastpt;
    int i=0,j=0,ctr=0;
    ctr=getbonds(mysite,bondpt);
    if(ctr==0)
    {
        //No cluster or monomer;
        bond=false;
    }
}

```

```

else
{
    //for each bond recieved
    bond=true;
    for(;j<ctr;j++)
    {
        pt=bondpt[j];
        //if monomer means it has an index
        if(neighborIsMonomer(pt,mysite))
        {
            //delete both you and monomer
            //delete monomer
            deletemonomer(pt);
        }
    }
    //delete your self IF you are a monomer
    if(mysite.index!=-1)
    {
        //rearrange list if mcount is greater than 1
        deletemonomer(mysite.pos);
    }
}
return bond;
}

void Lattice::upnbhd(site mysite)
{
    bool bond=false;
    point pt, bondpt[dir],lastpt;
    int x,y;
    int i=0,j=0,ctr=0;
    ctr=getnbhrs(mysite,bondpt);
    //cout<<"getcte="<<ctr<<endl;
    checksite(mysite);
    for(i=0;i<ctr;i++)
    {
        x=bondpt[i].x;
        y=bondpt[i].y;

        checksite(location[x][y]); //bond=checkupdatebonds(location[bondpt[i].x][bondpt[i].y]);
    }
    if(mysite.h<=0 && mysite.index!=-1)
    {
        deletemonomer(mysite.pos);
    }
}

void Lattice::checksite(site mysite)

```



```

{
    bool bond=false;
    bond=checkupdatebonds(location[mysite.pos.x][mysite.pos.y]);
    if(bond==false)
    {
        if((location[mysite.pos.x][mysite.pos.y].index== -1)
        &&(location[mysite.pos.x][mysite.pos.y].h>0))
        {
            location[mysite.pos.x][mysite.pos.y].index=mcount;
            monomerloc[mcount]=location[mysite.pos.x][mysite.pos.y].pos;
            mcount++;
            addmonomerchange(0,mysite.pos);
        }
    }
}

void Lattice::addmonomerchange(int tag,int lastpoint)
{
    change[changecount].time=time;
    if(tag==0)
    {
        myevents[eventcount].oldval=monomerloc[count];
        change[changecount].newsite=mcount;
        change[changecount].newval=monomerloc[mcount];
        change[changecount].tag=0;
        changecount++;
    }
    else
    {
        change[changecount].oldsite=lastpoint;
        change[changecount].oldval=monomerloc[lastpoint];
        change[changecount].newsite=count-1;
        change[changecount].newval=monomerloc[count-1];
        change[changecount].tag=1;
        changecount++;
    }
}

void Lattice::restorelist(float Ctime)
{
    int oldloc,newloc,tag;
    site oldval,newval;
    int j=changecount-1;
    while(change[j].time>Ctime)
    {
        oldloc=change[j].oldsite;
        newloc=change[j].newsite;
        oldval=change[j].oldval;
        newval=change[j].newval;
        tag=change[j].tag;

```

```

    if(tag==1)
    {
        monomerloc[newloc]=newval;
        monomerloc[oldloc]=oldval;
        mcount++;
    }
    else
    {
        monomerloc[newloc]=oldval;
        mcount--;
    }
    j--;
}
}
void Lattice::saveconfig()
{
    int j;
    for (j=0;j<mcount;j++)
    {
        oldlist.monomerloc[j]=monomerloc[j];
    }
    oldlist.mcount=mcount;
    oldlist.ndep=ndep;
}

void Lattice::restorelist()
{
    int j;
    for (j=0;j<oldlist.mcount;j++)
    {
        monomerloc[j]=oldlist.monomerloc[j];
    }
    mcount=oldlist.mcount;
    ndep=oldlist.ndep;
}

void Lattice::restoreLattice()
{
    undoevent();
    restorelist();
    int i,j;
    int x,y;
    /**clear lattice**/
    for(i=0;i<size+2;i++)
    {
        for(j=0;j<size;j++)
        {
            location[i][j].index=-1;
        }
    }
}

```

```

    }
    /** restore indexes*/
    for(i=0;i<mcount;i++)
    {
        x=monomerloc[i].x;
        y=monomerloc[i].y;
        location[x][y].index=i;
    }
}

void Lattice::restoreLattice(float Ctime)
{
    undoevent(Ctime);
    restorelist(Ctime);
    int i,j;
    int x,y;
    /**clear lattice**/
    for(i=0;i<size+2;i++)
    {
        for(j=0;j<size;j++)
        {
            location[i][j].index=-1;
        }
    }
    /** restore indexes*/
    for(i=0;i<mcount;i++)
    {
        x=monomerloc[i].x;
        y=monomerloc[i].y;
        location[x][y].index=i;
    }
}

void Lattice::addbdyevent(site oldsite,site newsite,float,int tag)
{
    //add boundary events to list
    int bdyrightcount=bdycount[right],bdyleftcount=bdycount[left];
    //erase pointers
    oldsite.index=-1;newsite.index=-1;
    if((oldsite.pos.x>=size) || (newsite.pos.x>=size))
    {
        if(oldsite.pos.x==size)
        {
            oldsite.pos.x=0;
        }
        if(oldsite.pos.x==size+1)
        {
            oldsite.pos.x=1;
        }
    }
}

```

```

        if(newsite.pos.x==size)
        {
            newsite.pos.x=0;
        }
        if(newsite.pos.x==size+1)
        {
            newsite.pos.x=1;
        }
        bdyevent[right][bdyrightcount].oldsite=oldsite;
        bdyevent[right][bdyrightcount].newsite=newsite;
        bdyevent[right][bdyrightcount].t=time;
        bdyevent[right][bdyrightcount].tag=tag;
        bdycount[right]++;
    }

    if((oldsite.pos.x<=1) || (newsite.pos.x<=1))
    {
        if(oldsite.pos.x==0)
        {
            oldsite.pos.x=size;
        }
        if(oldsite.pos.x==1)
        {
            oldsite.pos.x=size+1;
        }
        if(newsite.pos.x==0)
        {
            newsite.pos.x=size;
        }
        if(newsite.pos.x==1)
        {
            newsite.pos.x=size+1;
        }
        bdyevent[left][bdyleftcount].oldsite=oldsite;
        bdyevent[left][bdyleftcount].newsite=newsite;
        bdyevent[left][bdyleftcount].t=time;
        bdyevent[left][bdyleftcount].tag=tag;
        bdycount[left]++;
    }
}

void Lattice::deposit()
{
    /*
    1. Find a location
    2. place monomer in monomer list IF NO neighbours around!
    3. do not deposit on ghost region;
    */
    float xrand=ranlist[iran];

```

```

iran++;
float yrand=ranlist[iran];
iran++;

int locx=xrand*(size)+1;
int locy=yrand*(size);

//add height
location[locx][locy].h+=1;
if((locx==1) || (locy==size))
{
    //add event to bdylist
    addbdyevent(location[locx][locy],location[locx][locy],time,depevent);
    upnbhd(location[locx][locy]);
}
//cout<<"newdeploc.x= "<<locx<<" newdeploc.y= "<<locy<<endl;
//add to eventlist;
myeventlist[nevent].oldsite=location[locx][locy];
myeventlist[nevent].newsite=location[locx][locy];
myeventlist[nevent].ranseq=nevent;
myeventlist[nevent].t=time;
myeventlist[nevent].tag=depevent;
}

void Lattice::diffuse()
{
    point newloc,oldloc,lastloc;
    bool bonded=false;
    float ranm=ranlist[iran];
    iran++;
    if(mcount>0)
    {
        int loc=(ranm)*(mcount-1);
        oldloc=monomerloc[loc];
        newloc=ranmove(location[oldloc.x][oldloc.y]);
        location[newloc.x][newloc.y].h+=1;
        location[oldloc.x][oldloc.y].h-=1;
        //cout<<" locx"<<oldloc.x<<" locy"<<oldloc.y<<endl;
        //cout<<"nlocx"<<newloc.x<<"nlocy"<<newloc.y<<endl;
        //cout<<"bdyevent"<<bdycount<<endl;
        if(location[oldloc.x][oldloc.y].h<0)
        {
            cout<<"fuuuuuuuuuuuuuuuuck"<<endl;
        }
        //boundary event
        if((newloc.x<1) || (newloc.x>size))
        {
            //add event to bdylist
            deletemonomer(oldloc);

```

```

addbdyevent(location[oldloc.x][oldloc.y],location[newloc.x][newloc.y],time,diffevent);
    }
    if(newloc.x==1 || newloc.x==size)
    {

addbdyevent(location[oldloc.x][oldloc.y],location[newloc.x][newloc.y],time,diffevent);
    }
    if((oldloc.x<=1 )|| (oldloc.x>=size))
    {
        //add event to bdylist

addbdyevent(location[oldloc.x][oldloc.y],location[newloc.x][newloc.y],time,diffevent);
    }
    //diffusion may release trapped monomer but capture
    released monomer
    if(location[oldloc.x][oldloc.y].h>location[newloc.x][newloc.y].h)
    {
        upnbhd(location[oldloc.x][oldloc.y]);    }
    else
    {
        //Move Monomer by changing index location
        location[newloc.x][newloc.y].index=location[oldloc.x][oldloc.y].index;

monomerloc[location[oldloc.x][oldloc.y].index]=location[newloc.x][newloc.y].pos;
        location[oldloc.x][oldloc.y].index=-1;
        upnbhd(location[newloc.x][newloc.y]);
    }
    }
    //add to eventlist;
    myeventlist[nevent].oldsite=location[oldloc.x][oldloc.y];
    myeventlist[nevent].newsite=location[newloc.x][newloc.y];
    myeventlist[nevent].ranseq=nevent;
    myeventlist[nevent].t=time;
    myeventlist[nevent].tag=diffevent;
    //cout<<"****End Diffuson*****"<<endl;
}
void Lattice::savebdylist()
{
    int a,b,i,bdyrightcrec=bdycountrec[right],bdyleftcrec=bdycountrec[left];

    for(i=0;i<bdyleftcrec;i++)
    {
        oldbdyeventrec[left][i]=bdyeventrec[left][i];
    }
    oldbdycountrec[left]=bdyleftcrec;

    for(i=0;i<bdyrightcrec;i++)

```

```

    {
        oldbdyeventrec[right][i]=bdyeventrec[right][i];
    }

    oldbdycountrec[right]=bdyrightcrec;
}

int Lattice::comparebdylist()
{
    int a, b, acheck, bcheck;

    acheck = 0;
    bcheck = 0;
    for (a=0; a < 2; a++) {
        if (oldbdycountrec[a]!= bdycountrec[a]) {
            redoflag = 1;
            acheck = 1;
        } else {
            for (b=0; b < bdycountrec[a]; ) {
                if (oldbdyeventrec[a][b].t != bdyeventrec[a][b].t) {
                    redoflag = 1;
                    bcheck = 1;
                    b = bdycountrec[a];
                }
                if (oldbdyeventrec[a][b].newsite.pos.x !=
bdyeventrec[a][b].newsite.pos.x) {
                    redoflag = 1;
                    bcheck = 1;
                    b = bdycountrec[a];
                }
                if
(oldbdyeventrec[a][b].newsite.pos.y!=bdyeventrec[a][b].newsite.pos.y) {
                    redoflag = 1;
                    bcheck = 1;
                    b = bdycountrec[a];
                }
                if (oldbdyeventrec[a][b].newsite.h!=bdyeventrec[a][b].newsite.h) {
                    redoflag = 1;
                    bcheck = 1;
                    b = bdycountrec[a];
                }
                b++;
            }
        }
    }
}

```

```

void Lattice::p()
{
    cout<<"*****S*****";
    float theta=0,vacancy=(float) mcount ,lat=(float)latsize;
    float x;
    for (int i=0;i<size;i++)
    {
        cout<<endl;
        for (int j=0;j<size+2;j++)
            cout<<location[j][i].h;
    }

    cout<<endl<<"mcount="<<mcount<<endl;

    //cout<<"*****E*****"<<endl;
    theta=(lat-vacancy)/lat;
    //cout<<"Theta="<<theta<<endl;
    for (int i=0;i<=size;i++)
    {
        cout<<endl;
        for (int j=0;j<size+2;j++)
            cout<<location[j][i].index<<" ";
    }
}

void Lattice::doKMC()
{
    ///create and save random number
    float ranX;

    ranX=ranlist[iran];
    iran++;

    float Trate,Drate;
    Drate=.25*mcount*difrate;
    Trate=Drate+(deprate* (float) latsize);

    float prob=(Drate/Trate);

    //cout<<"ranX="<<ranlist[iran]<<endl;
    if(ranX<prob)
        diffuse();
    else
    {
        deposit();
        ndep++;
    }
    nevent++;
}

```



```

void Lattice::undoevent() {
    int a, xi, yi, xf, yf, tag;
    double t;

    if (redoflag == 0) {
        return;
    }

    for (a=nevent-1; a >=0; a--) {
        tag = myeventlist[a].tag;
        xi = myeventlist[a].oldsite.pos.x;
        yi = myeventlist[a].oldsite.pos.y;
        xf = myeventlist[a].newsite.pos.x;
        yf = myeventlist[a].newsite.pos.y;
        t = myeventlist[a].t;

        switch (tag) {
            case 0:
                location[xi][yi].h = myeventlist[a].oldsite.h;
                if (myeventlist[a].newsite.h != -1)
                    location[xf][yf].h = myeventlist[a].newsite.h;
                break;
            case 1:
                location[xi][yi].h = location[xi][yi].h + 1;
                location[xf][yf].h = location[xf][yf].h - 1;
                break;
            case 2:
                location[xi][yi].h = location[xi][yi].h - 1;
                ndep--;
                break;
            default:
                cout<<"Error in tag"<<endl;
                return; /* SHOULDN'T THIS EXIT()?!!? */
        }
    }
}

void Lattice::undoevent(float Ctime) {
    int a, xi, yi, xf, yf, tag;
    double t;

    if (redoflag == 0) {
        return;
    }
    a=nevent-1;
    t = myeventlist[a].t;
    while(t>Ctime) {
        tag = myeventlist[a].tag;
        xi = myeventlist[a].oldsite.pos.x;

```

```

        yi  = myeventlist[a].oldsite.pos.y;
        xf  = myeventlist[a].newsite.pos.x;
        yf  = myeventlist[a].newsite.pos.y;
        t   = myeventlist[a].t;

        switch (tag) {
        case 0:
            location[xi][yi].h = myeventlist[a].oldsite.h;
            if (myeventlist[a].newsite.h != -1)
                location[xf][yf].h = myeventlist[a].newsite.h;
            break;
        case 1:
            location[xi][yi].h = location[xi][yi].h + 1;
            location[xf][yf].h = location[xf][yf].h - 1;
            break;
        case 2:
            location[xi][yi].h = location[xi][yi].h - 1;
            ndep--;
            break;
            default:
                cout<<"Error in tag"<<endl;
                return; /* SHOULDN'T THIS EXIT()?? */
        }
        a--;
    }
}

void Lattice::randgen()
{
    int i;
    for(i=0;i<10000;i++)
    {
        ranlist[i]=((float)rand()/(float)RAND_MAX);
    }
}

void Lattice::updateBuffer(int iranflag) {
    int a, b, aml, x, y, xi, ii, abflag, mflag, sdir, dir, aid, i, j, hij,
    hxy, tag;
    double newTrate, oldTrate;
    point oldsite, newsite;
    i=sortbdyevent[nupdate].oldsite.pos.x;
    j=sortbdyevent[nupdate].oldsite.pos.y;
    x=sortbdyevent[nupdate].newsite.pos.x;
    y=sortbdyevent[nupdate].newsite.pos.y;
    tag=sortbdyevent[nupdate].tag;

    time = sortbdyevent[nupdate].t;
    if (redoflag == 0) {
        return;
    }
}

```

```

    }
    //cout<<"update buffer!"<<endl;
    //cout<<"x="<<x<<" y="<<nupdate<<endl;
    //cout<<"i="<<x<<" j="<<y<<endl;
        if(tag==diffevent)
    {
        location[i][j].h = sortbdyevent[nupdate].oldsite.h;
        upnbhd(sortbdyevent[nupdate].oldsite);
    }
    else
    {
        location[x][y].h = sortbdyevent[nupdate].newsite.h;
        upnbhd(sortbdyevent[nupdate].newsite);
    }
    /* add this event in my event list */
    myeventlist[nevent].oldsite=location[x][y];
    myeventlist[nevent].newsite=location[i][j];
    myeventlist[nevent].ranseq=iran - iranflag;
    myeventlist[nevent].t=time;
    myeventlist[nevent].tag=0;
    nupdate++;
    nevent++;
}

void Lattice::sorting_nbevent() {
    int a, b, nxcv, i, j, caselabel, dir, idn;
    double t;
    boundaryevent swap;

    if (bdycountrec[left] > 0 && bdycountrec[right] == 0)
        caselabel = 0;
    if (bdycountrec[left] == 0 && bdycountrec[right] > 0)
        caselabel = 1;
    if (bdycountrec[left] > 0 && bdycountrec[right] > 0)
        caselabel = 2;

    switch(caselabel) {
    case 0:
        for (a=0; a < bdycountrec[left]; a++) {
            sortbdyevent[a]= bdyeventrec[0][a];
        }
        tnbdyevent = bdycountrec[0];
        break;
    case 1:
        for (a=0; a < bdycountrec[1]; a++) {
            sortbdyevent[a] = bdyeventrec[1][a];
        }
        tnbdyevent = bdycountrec[1];
        break;
    }

```

```

case 2:
    tnbdyevent = bdycountrec[0] + bdycountrec[1];
    nxcv = 0;
    /* sort the events in early time order */
    for (a = 0; a < bdycountrec[0]; a++) {
        sortbdyevent[nxcv] = bdyeventrec[0][a];
        nxcv++;
    }

    for (a=0; a < bdycountrec[1]; a++) {
        sortbdyevent[nxcv] = bdyeventrec[1][a];
        nxcv++;
    }

    for (j=1; j < tnbdyevent; j++) {
        swap=sortbdyevent[j];
        i = j - 1;
        while (i >= 0 && sortbdyevent[i].t > t) {
            sortbdyevent[i+1] = sortbdyevent[i];
            i--;
        }
        sortbdyevent[i+1]=swap;
    }
    break;
default:
    break;
}
}

```

A.4 synch.cpp

Code: synch.cpp

```

#include "lattice.h"
extern MPIWrapper mpi;
void synch(Lattice * newlatt)
{
    int ctr,iranflag;
    int tchange=1;

```

```

float tmytime;
newlatt->subcycle=1;
int ctcycles=0;
while(tchange>0 && ctcycles<20)
{
    ctcycles++;
    if(newlatt->myid==1)
    {
        // cout<<"id="<<newlatt->myid<<"count
cycles"<<ctcycles<<endl;
    }

    /* start iteration from here */
    //newlatt->undoflag = -1;
    newlatt->redoflag = 0;
    newlatt->tnbdyevent = 0;

    /* check whether new iteration is needed */
    if (newlatt->subcycle == 1) {
        if (newlatt->bdycountrec[0] + newlatt->bdycountrec[1] > 0) {
            newlatt->redoflag = 1;
            newlatt->sorting_nbevent();
        }
    } else {
        newlatt->comparebdylist();
        if (newlatt->redoflag == 1) {
            newlatt->savebdylist();
            if (newlatt->bdycountrec[0] +
newlatt->bdycountrec[1] > 0)
                newlatt->sorting_nbevent();
        }
    }

    /* new iteration is needed: newlatt->redoflag=1 */
    if (newlatt->redoflag == 1) {
        newlatt->restoreLattice(); /* restore starting
configuration */

        newlatt->nupdate = 0;
        newlatt->time = 0.0;
        newlatt->iran = 0;
        newlatt->nevent = 0;

        newlatt->myeventlist[newlatt->nevent].ranseq = newlatt->iran;
        newlatt->calctime();

        /* save numbers of changes */
        /* repeat kmc event : update buffers and start from
there*/

```

```

        /* newlatt->time is later than 1st boundary event
time */
        if (newlatt->time > newlatt->T && newlatt->tnbdyevent>0) {
            tmytime = newlatt->T+1.0;
            for (; tmytime > newlatt->T && newlatt->nupdate <
newlatt->tnbdyevent;) {
                iranflag = 0;
                newlatt->updateBuffer(iranflag);
                newlatt->myeventlist[newlatt->nevent].ranseq =
newlatt->iran;
                newlatt->calctime();
                tmytime = newlatt->time;
            }
        }

        /* newlatt->time is earlier than 1st boundary event
time */
        while (newlatt->time < newlatt->T) {
            if (newlatt->time < newlatt->T) {
                if (newlatt->nupdate < newlatt->tnbdyevent) {
                    if (newlatt->time <
newlatt->sortbdyevent[newlatt->nupdate].t) {
                        newlatt->doKMC();
                    } else {
                        iranflag = 1;
                        newlatt->updateBuffer(iranflag);
                    }
                } else {
                    newlatt->doKMC();
                }
            }
            newlatt->myeventlist[newlatt->nevent].ranseq =
newlatt->iran;
            newlatt->calctime();
            tmytime = newlatt->time;
            for (; tmytime > newlatt->T && newlatt->nupdate <
newlatt->tnbdyevent;) {
                iranflag = 0;
                newlatt->updateBuffer(iranflag);
                newlatt->myeventlist[newlatt->nevent].ranseq =
newlatt->iran;
                newlatt->calctime();
                tmytime = newlatt->time;
            }
            //cout<<"I am"<<newlatt->time<<endl;

            //cout<<"subcycle*****"<<newlatt->redoflag<<"flag*****"
"<<ctr<<endl;
        }

```

```

    }
    /* check how many processors have a change in
the previous events */
    newlatt->subcycle++; /* increase number of iteration */

    /* check how many processors were redone */
    mpi.allReduce(&newlatt->redoflag, &tchange, 1, MPI_INT, MPI_SUM);

    if (tchange > 0) { /* some processors are unhappy: redo must
be needed */
        newlatt->bdycountrec[left]=0;
        newlatt->bdycountrec[right]=0;
        sendmsgs(newlatt);
    }
}

```

A.5 comm.cpp

Code: comm.cpp

```

#include "lattice.h"
extern MPIWrapper mpi;
void sendmsgs(Lattice * newlatt)
{

    mpi.sendboundaryevent(newlatt->bdyevent[left], newlatt->bdycount[left], newlatt->nbhr[left]);

    mpi.sendboundaryevent(newlatt->bdyevent[right], newlatt->bdycount[right], newlatt->nbhr[right]);

    mpi.recvboundaryevent(newlatt->bdyeventrec[left], newlatt->nbhr[left]);
    mpi.recvboundaryevent(newlatt->bdyeventrec[right], newlatt->nbhr[right]);

    for(int j=0; j<2; j++)
    {
        newlatt->bdycount[j]=0;
    }
}

```

```
}  
}
```

Appendix B

Time Warp Code

B.1 main.cpp

Code: main.cpp

```
#include <iostream>
using std::cout;
using std::cerr;
using std::endl;

#include <fstream>
using std::ofstream;

#include "exception.h"
#include "lattice.h"
#include "mpiwrapper.h"

const int globalSyncThreshold = 1000;

string makeFileName(string, string, int);

int main(int argc, char* argv[]) {
    Lattice lattice;
    string pngFilename = "";
    string logFilename = "";
    fstream logFile;
    double minGlobalTime = 0.0;
    double maxGlobalTime = 0.0;
    int globalTimeCounter = 0;
    double gConvergence = 0.0;
    int eventCount = 0;
```

```

// setup the lattice mpi stuff
lattice.mpi.init(&argc,&argv);

try {

    pngFilename = makeFileName("height-node","png",lattice.mpi.getRank());
    logFilename = makeFileName("log","txt",lattice.mpi.getRank());

    logFile.open(logFilename.c_str(),fstream::out|fstream::trunc);

    if(!logFile) {
        string error = "Couldn't open log file " + logFilename;
        throw(Exception(error));
    }

    lattice.setMinGlobalTime(0.0);

    lattice.mpi.barrier();

    // MAIN LOOP
    while(gConvergence < 1.0) {
        // retrieve any remote events
        lattice.negotiateEvents(logFile);

        // do the next event
        lattice.doNextEvent();
        // see if it's time for a global sync
        if(globalTimeCounter > globalSyncThreshold) {
            lattice.mpi.barrier();

            lattice.negotiateEvents(logFile);

            // allreduce to find the min time
            minGlobalTime =
lattice.mpi.allReduceDouble(lattice.getLocalTime(),MPI_MIN);
            maxGlobalTime =
lattice.mpi.allReduceDouble(lattice.getLocalTime(),MPI_MAX);
            eventCount =
lattice.mpi.allReduceInt(lattice.getEventCount(),MPI_SUM);

            // set the global time in the lattice
            lattice.setMinGlobalTime(minGlobalTime);

            // clear the counter
            globalTimeCounter = 0;

            // calculate the global convergence
            gConvergence = (double)eventCount/(double)(lattice.mpi.getNodeCount())

```

```

* SIZE);

    if(lattice.mpi.isRoot()) {
        cout << minGlobalTime << " " << maxGlobalTime << " " << gConvergence
<< endl;
        cout.flush();
    }
} else
    ++globalTimeCounter;
}

logFile << "exit main loop" << endl;
logFile.flush();

lattice.mpi.barrier();

// rollback to minimum global time

//lattice.printLatticeHeight(logFile);
logFile << "gCovergence = " << gConvergence << endl;
lattice.printStats(logFile);
lattice.createHeightMap(pngFilename);
lattice.mpi.printStats(logFile);

lattice.cleanup(logFile);

logFile.close();

lattice.mpi.barrier();

} catch(Exception err) {
    cerr << err.error << endl;
}

// close the mpi stuff
lattice.mpi.shutdown();
return(0);
}

string makeFileName(string prefix, string ext, int rank) {
    string output = prefix + ".";
    output += (char)('a' + rank);
    return(output + "." + ext);
}

```

B.2 lattice.h

Code: lattice.h

```
#include <vector>
using std::vector;

#include <queue>
using std::priority_queue;

#include <stack>
using std::stack;

#include <iomanip>
using std::setw;
using std::hex;
using std::dec;
using std::setprecision;

#include <fstream>
using std::fstream;

#include <string>
using std::string;

#include <png.h>

#include "exception.h"
#include "latprim.h"
#include "latconst.h"
#include "event.h"
#include "randgen.h"
#include "rewindlist.h"
#include "mpiwrapper.h"

#ifndef LATTICE_H
#define LATTICE_H

#define GET_DIR(a) ((a < LEFT_X_BOUNDARY) ? LEFT : RIGHT)

class Lattice {
public:
    Lattice();
```

```

~ Lattice();

void cleanup(fstream&);

bool doNextEvent();

double getLocalTime() {
    return(localTime);
}

bool setMinGlobalTime(double mGT) {
    minGlobalTime = mGT;
    return(true);
}

double getMinGlobalTime() {
    return(minGlobalTime);
}

bool negotiateEvents(fstream&);

// DEBUG FUNCTIONS
void printLatticeHeight(fstream& file) {
    for(int i=0; i < DIM_X + GHOST + GHOST; ++i) {
        for(int j=0; j < DIM_Y; ++j) {
            file << lattice[i][j].h << " ";
        }
        file << endl;
    }
    file << "-----" << endl <<
endl;
    file.flush();
}

void printLatticeIndex(fstream& file) {
    for(int i=0; i < DIM_X + GHOST + GHOST; ++i) {
        for(int j=0; j < DIM_Y; ++j) {
            if(lattice[i][j].listIndex >= 0)
                file << setw(4) << lattice[i][j].listIndex;
            else
                file << setw(4) << "x";
            file << " ";
        }
        file << endl;
    }
    file << "-----" << endl <<
endl;
    file.flush();
}

```

```

void printMonomerList(fstream& file) {
    file << "monomerList[" << monomerList.size() << "]" at time=" << localTime
<< endl;
    for(int i=0; i < monomerList.size(); ++i) {
        site* s = monomerList[i];
        file << i << ": (" << s->p.x << ", " << s->p.y << ") h=" << s->h << "
listIndex=" << s->listIndex << " " << hex << s << dec << endl;
    }
    file << "-----" << endl <<
endl;
    file.flush();
}

void printStats(fstream& file) {
    file << setprecision(10) << endl;
    file << "COLLECTED STATISTICS" << endl;
    file << "-----" << endl;
    file << "Convergence = " << getConvergence() << endl;
    file << "Total Event Count = " << countEvents << endl;
    file << "Total Diffusion Events = " << countDiffusion << endl;
    file << "Total Deposition Events = " << (countEvents - countDiffusion) <<
endl;
    file << "Total Boundry Events = " << countBoundry << endl;
    file << "Total Number Remote Events = " << countRemote << endl;
    file << "Total Rollbacks Performed = " << countRollback << endl;
    file << "Monomer List Size = " << monomerList.size() << endl;
    file << "Local Time = " << localTime << endl;
    file << "Min Global Time = " << minGlobalTime << endl;
    file << "Size = " << SIZE << " DIM_X = " << DIM_X << " DIM_Y = " << DIM_Y
<< endl;
    file << "-----" << endl <<
endl;
    file.flush();
}

int getEventCount() { return(countEvents - countDiffusion); }

bool createHeightMap(string filename);

MPIWrapper mpi; // I can't think of a better place for this.

bool rollback(const double);

double getConvergence() {
    return((double)(countEvents - countDiffusion) / (double)(DIM_X * DIM_Y));
}

private:

```

```

double computeTime();
bool deposit();
bool diffuse();
bool doKMC();
EventType getNextEventType();
bool commitEvent(Event*);
site* randomMove(site*);
bool isBoundry(point);
bool isBound(site*);
bool clearBonded(site*,const double);
bool translateMessages(vector<Event*>* , vector<message*>);
message* makeMessage(Event*);
bool hasAntiEvent(Event*);

double localTime;    // the time local to the lattice
double minGlobalTime; // the minimum Global time (point of no
return)

RewindList<site *> monomerList; // list of all unbound monomers
//site lattice[DIM_X + GHOST + GHOST][DIM_Y]; // the lattice
(the extra two are the ghost region)
site** lattice;

float depositionRate; // the deposition rate of monomers
float diffusionRate;  // the diffusion rate of monomers

int countDiffusion;
int countEvents;
int countBoundry;
int countRemote;
int countRollback;

priority_queue<Event*> remoteEventList; // list of all the remote
dep/diffusion events
stack<Event*> eventList;                // stack of all events to
rollback the simulation
vector<Event*> antiEvents;              // list of anti-events that
will occur in the future

RandGen rng; // random number generator

point movementDir[NUM_DIR]; // array of movement types
message m; // message for sending events
};

#endif

```

B.3 latconst.h

Code: latconst.h

```
#ifndef LATCONST_H
#define LATCONST_H

// lattice dimensions and area
//const int DIM_X = 4096;
//const int DIM_X = 2048;
const int DIM_X = 1024;           // the total X dimension of
the lattice                      // the total X
//const int DIM_X = 512;           // the total X
dimension of the lattice         // the total Y
//const int DIM_Y = 1024;         // the total Y
dimension of the lattice
const int DIM_Y = 256;
const int GHOST = 1;             // the size of the ghost
region
const int LEFT_X_BOUNDARY = GHOST; // the left boundary
const int RIGHT_X_BOUNDARY = DIM_X - GHOST - 1; // the right boundary
const int SIZE = DIM_X * DIM_Y;   // the area (number) of
sites in the lattice

// enviroment attributes
const int NUM_DIR = 8; // number of directions

// minimum local time
const double MINIMUM_TIME = 0.001;

// random number vars
const int NUMBER_START = 10000; // the number of random numbers to
start with
const int SEED = 0;              // the initial seed

#endif
```

B.4 latprim.h

Code: latprim.h

```
#ifndef LATPRIM_H
#define LATPRIM_H

typedef struct {
    int x;
    int y;
} point;

typedef struct {
    point p;
    int h;
    int listIndex;
} site;

typedef point* pointPTR; // pointer to a point
typedef site* sitePTR;   // pointer to a site

#endif
```

B.5 lattice.cpp

Code: lattice.cpp

```
#include "lattice.h"

Lattice::Lattice() : localTime(MINIMUM_TIME), minGlobalTime(0.0),
    depositionRate(1.0), diffusionRate(1.0e6),
```

```

        countDiffusion(0), countEvents(0), countBoundry(0), countRemote(0),
countRollback(0), rng(10000,7) {

    lattice = new site*[DIM_X+GHOST+GHOST];

    // set up the lattice array
    for(int i=0; i < DIM_X + GHOST + GHOST; i++) {
        lattice[i] = new site[DIM_Y];
        for(int j=0; j < DIM_Y; j++) {
            lattice[i][j].p.x = i;
            lattice[i][j].p.y = j;
            lattice[i][j].h = 0;
            lattice[i][j].listIndex = -1;
        }
    }

    // set the directions
    movementDir[0].y= 1; movementDir[0].x= 0;
    movementDir[1].y= 0; movementDir[1].x= 1;
    movementDir[2].y=-1; movementDir[2].x= 0;
    movementDir[3].y= 0; movementDir[3].x=-1;
    movementDir[4].y= 1; movementDir[4].x= 1;
    movementDir[5].y= 1; movementDir[5].x=-1;
    movementDir[6].y=-1; movementDir[6].x= 1;
    movementDir[7].y=-1; movementDir[7].x=-1;
}

Lattice::~Lattice() {
    for(int i=0; i < DIM_X + GHOST + GHOST; ++i) {
        delete [] lattice[i];
    }
    delete [] lattice;
}

void Lattice::cleanup(fstream& logFile) {
    vector<message> m;

    logFile << ((mpi.isMessage()) ? "messages waiting" : "message queue empty")
<< endl;
    logFile << ((mpi.isAntiMessage()) ? "messages waiting" : "message queue
empty") << endl;
    if(mpi.isMessage()) {
        mpi.recvMessages(&m);
        m.empty();
    }
    if(mpi.isAntiMessage()) {
        mpi.recvAntiMessages(&m);
        m.empty();
    }
}
}

```

```

bool Lattice::doNextEvent() {
    double nextKMCTime = computeTime();
    EventType nextEventType;

    if(remoteEventList.empty() || nextKMCTime < remoteEventList.top()->time) {
        // set the local time to the nextKMCTime
        localTime = nextKMCTime;

        // get the event type
        nextEventType = getNextEventType();

        // depending on the event type commit different events
        switch(nextEventType) {
            case eventDeposition: {
                point p;

                // set up the new site object
                p.x = (int)(rng.getRandom(localTime) * (DIM_X));
                p.y = (int)(rng.getRandom(localTime) * DIM_Y);

                // commit the event
                commitEvent(new
Event(&lattice[p.x][p.y],localTime,true,nextEventType));
            } break;
            case eventDiffusion: {
                site *oldSite, *newSite;
                int index = (int)(rng.getRandom(localTime) * monomerList.size());

                // make sure there actually a monomer in the system to
diffuse
                if(monomerList.size() > 0) {
                    // set the old site (random monomer from the list) and
get a new site
                    oldSite = monomerList[index];
                    newSite = randomMove(oldSite);
                    monomerList.add(newSite,localTime);
                    monomerList.remove(index,localTime);

                    // set the index of the new site and update our entry in
the monomerList
                    newSite->listIndex = index;
                    monomerList[index] = newSite;

                    // clear the listIndex of the old site
                    oldSite->listIndex = -1;

                    // commit the event

```

```

        commitEvent(new
Event(oldSite,newSite,localTime,true,nextEventType,index));

    }
    } break;
}
} else {
    // see if there is a waiting anti-event for this remote event
    if(hasAntiEvent(remoteEventList.top())) {
        // delete the remote event to play nice with memory
        delete remoteEventList.top();

        // remove the remote event from the queue
        remoteEventList.pop();

        // return false abandoning trying to do an event this cycle
        return(false);
    }

    // set nextKMCTime to the remote time
    localTime = remoteEventList.top()->time;

    // commit the remote event to the simulation
    commitEvent(remoteEventList.top());

    // remove the remote event from the queue
    remoteEventList.pop();

    countRemote++;
}

// increment the event counter
++countEvents;

return(true);
}

double Lattice::computeTime() {
    double Drate = diffusionRate * monomerList.size() * 0.25f;
    double totaldep = depositionRate * SIZE;
    double dt = -log(rng.getRandom(localTime))/(Drate+totaldep);

    return(localTime + dt);
}

EventType Lattice::getNextEventType() {
    float Drate = 0.25f * monomerList.size() * diffusionRate;
    float Trate = Drate + (depositionRate * SIZE);
    float probab = (Drate / Trate);

```

```

    // if the next random number from the stream is less then the
    probability
    // the the next event is a eventDiffusion, return a diffusion
    event
    if(rng.getRandom(localTime) < prob)
        return(eventDiffusion);

    // catch-all is the deposition event
    return(eventDeposition);
}

bool Lattice::commitEvent(Event* event) {

    if(event == NULL)
        throw(Exception("NULL event passed to Lattice::commitEvent()!"));

    // process the event based on the event type
    switch(event->eventType) {
    case eventDeposition: {
        /// increment the height up
        ++(event->newSite->h);

        // see if the monomer falls on the boundry (x ==
        LEFT_X_BOUNDARY || x == RIGHT_X_BOUNDARY)
        if(isBoundry(event->newSite->p) && event->isLocal) {
            // send the event off to the correct neighbor
            mpi.sendMessage(makeMessage(event), GET_DIR(event->newSite->p.x));

            // set the boundry event flag
            event->isBoundry = true;

            // increment the countBoundry
            ++countBoundry;
        }

        // see if the monomer will bond
        if(!isBound(event->newSite)) {
            // no bond was formed so move on
            event->newSite->listIndex =
            monomerList.add(event->newSite, event->time);
        }

    } break;
    case eventDiffusion: {
        // increment the new site height up
        ++(event->newSite->h);

        // increment the old site height down

```

```

--(event->oldSite->h);

// see if the monomers new site falls on the boundry (x == 1
// x == DIM - 1)
if(isBoundry(event->newSite->p) && event->isLocal) {
    // send the event off to the correct neighbor
    mpi.sendMessage(makeMessage(event),GET_DIR(event->newSite->p.x));

    // set the boundry event flag
    event->isBoundry = true;

    // increment the countBoundry
    ++countBoundry;
}

// see if the monomer will bond
if(isBound(event->newSite)) {
    // remove the current monomer from the monomer list
    site* s = monomerList.remove(event->newSite->listIndex,event->time);

    // because of the way RewindList works remove() returns the
value of
    // the new element at the position, and we'll need to
change it's
    // listIndex value to reflect it's new index
    if(s != NULL)
        s->listIndex = event->newSite->listIndex;

    // invalidate the index of the newSite
    event->newSite->listIndex = -1;

    // clear any neighbors that could be unbound monomers from
the list
    clearBonded(event->newSite,event->time);
}

// increment the diffusion counter
++countDiffusion;
} break;
default: throw(Exception("commitEvent(): Invalid Event type!"));
}

// push the event into the event list
eventList.push(event);

return(true);
}

site* Lattice::randomMove(site* oldSite) {

```

```

    point newPoint;
    int i = (int)(rng.getRandom(localTime) * 4);

    // randomly move in a random direction with help from our
    direction array
    newPoint.x = oldSite->p.x + movementDir[i].x;
    newPoint.y = oldSite->p.y + movementDir[i].y;

    // make sure we don't leave our domain
    if(newPoint.x >= RIGHT_X_BOUNDARY)
        newPoint.x -= 2;
    if(newPoint.x <= LEFT_X_BOUNDARY)
        newPoint.x += 2;

    if(newPoint.y >= DIM_Y)
        newPoint.y -= 2;
    if(newPoint.y < 0)
        newPoint.y += 2;

    // return the new site on the lattice (where the monomer moved
    to)
    return(&lattice[newPoint.x][newPoint.y]);
}

bool Lattice::isBoundry(point p) {
    // if p.x falls in the ghost (0, DIM-1) or boundry (1, DIM-2)
    return true
    if(p.x < LEFT_X_BOUNDARY || p.x > RIGHT_X_BOUNDARY)
        return(true);

    // default is false
    return(false);
}

bool Lattice::isBound(site* s) {
    point p;

    // loop and check each direction, returning true if a neighbor
    is at the
    // same height or higher
    for(int i=0; i < NUM_DIR; ++i) {
        p = s->p;
        p.x += movementDir[i].x;
        p.y += movementDir[i].y;

        if(p.x >= 0 && p.x < DIM_X+GHOST+GHOST && p.y >= 0 && p.y < DIM_Y)
            if(lattice[p.x][p.y].h >= s->h)
                return(true);
    }
}

```

```

    // default is false
    return(false);
}

bool Lattice::clearBonded(site* s, const double t) {
    point p;

    // loop and inspect the neighbors marking any bonded and
    for(int i=0; i < NUM_DIR; ++i) {
        p = s->p;
        p.x += movementDir[i].x;
        p.y += movementDir[i].y;

        // if the point is a valid point and if the point is going to
        bound,
        // and if the listIndex is valid then delete it from the
        monomer list
        if(p.x >= RIGHT_X_BOUNDARY && p.x < RIGHT_X_BOUNDARY && p.y >= 0 && p.y <
        DIM_Y)
            if(lattice[p.x][p.y].h >= s->h)
                if(lattice[p.x][p.y].listIndex != -1) {
                    monomerList.remove(lattice[p.x][p.y].listIndex,t);
                    lattice[p.x][p.y].listIndex = -1;
                }
    }

    return(true);
}

bool Lattice::createHeightMap(string filename) {
    int x, y;
    int width=DIM_X, height=DIM_Y;
    png_byte color_type=PNG_COLOR_TYPE_RGBA;
    png_byte bit_depth=8;
    png_structp png_ptr;
    png_infop info_ptr;
    int number_of_passes=1;
    png_bytep * row_pointers;
    FILE* fp;

    row_pointers = (png_bytep*) malloc(sizeof(png_bytep) * height);
    for (y=0; y<height; y++)
        row_pointers[y] = (png_byte*) malloc(width*((bit_depth/8)*4));

    for (y=0; y<height; y++) {
        png_byte* row = row_pointers[y];
        for (x=0; x<width; x++) {
            png_byte* ptr = &(row[x*4]);

```



```

        if(lattice[x][y].h > 0) {
            ptr[0] = 255 - ((lattice[x][y].h*10)%255); ptr[1] = 255 -
((lattice[x][y].h*10)%255); ptr[2] = 255 - ((lattice[x][y].h*10)%255); ptr[3]
= 255;
        } else {
            ptr[0] = 0; ptr[1] = 0; ptr[2] = 0; ptr[3] = 255;
        }
    }
}

fp = fopen(filename.c_str(), "wb");
if(fp == NULL) {
    throw Exception("Couldn't open height map file!");
    return(false);
} else {
    png_ptr = png_create_write_struct(PNG_LIBPNG_VER_STRING, NULL, NULL,
NULL);
    info_ptr = png_create_info_struct(png_ptr);
    png_init_io(png_ptr, fp);
    png_set_IHDR(png_ptr, info_ptr, width, height,
        bit_depth, color_type, PNG_INTERLACE_NONE,
        PNG_COMPRESSION_TYPE_BASE, PNG_FILTER_TYPE_BASE);
    png_write_info(png_ptr, info_ptr);
    png_write_image(png_ptr, row_pointers);
    png_write_end(png_ptr, NULL);
}

return(true);
}

bool Lattice::rollback(double t) {

    if(t < minGlobalTime)
        t = minGlobalTime;

    // rewind the lattice using the eventList stack
    while(!eventList.empty() && localTime > t) {
        Event* event = eventList.top();

        switch(event->eventType) {
            case eventDeposition: {
                // decrease the height of the deposition site
                --event->newSite->h;

                // clear the listIndex just in case
                event->newSite->listIndex = -1;

                // see if the monomers new site falls on the boundry (x ==
1 || x == DIM - 1)

```

```

        if(isBoundry(event->newSite->p)) {
            if(event->isLocal) {
                // send the anti-message

mpi.sendAntiMessage(makeMessage(event),GET_DIR(event->newSite->p.x));
            } else {
                // store the event back in the remoteEvents list
                remoteEventList.push(event);
                countRemote--;
            }

            // decrement the countBoundry
            --countBoundry;
        }

    }break;
case eventDiffusion: {
    // decrease the height of the new site
    --event->newSite->h;

    // increase the height of the old site
    ++event->oldSite->h;

    // clear the listIndex just in case
    event->newSite->listIndex = -1;
    event->oldSite->listIndex = -1;

    // see if the monomers new site falls on the boundry (x ==
1 // x == DIM - 1)
    if(isBoundry(event->newSite->p)) {
        if(event->isLocal) {
            // send the anti-message

mpi.sendAntiMessage(makeMessage(event),GET_DIR(event->newSite->p.x));
        } else {
            // store the event back in the remoteEvents list
            remoteEventList.push(event);
            countRemote--;
        }

        // decrement the countBoundry
        --countBoundry;
    }

    // decrement the countDiffusion stat
    --countDiffusion;
}break;
}

```

```

    // set the local clock to the event time
    localTime = event->time;

    // decrement the countEvents;
    --countEvents;

    // if it's a local event clean up the memory
    if(event->isLocal)
        delete event;

    // pop the top of the event list
    eventList.pop();
}

// rewind the RNG
rng.rewind(t);

// rewind the monomerList
monomerList.rewind(t);

// off by one error
++countEvents;

// fix all of the listIndex entries
for(int i=0; i < monomerList.size(); ++i) {
    monomerList[i]->listIndex = i;
}

++countRollback;

return(true);
}

bool Lattice::negoitateEvents(fstream& logFile) {
    vector<message> messages;
    vector<message> antiMessages;
    vector<Event*> remoteEvents;
    float pastTime = 0.0;
    bool isRollback = false;

    // get any waiting message
    // logFile << "lattice.mpi.recvMessages()" << endl;
    logFile.flush();
    mpi.recvMessages(&messages);

    // get any waiting anti-messages
    // logFile << "lattice.mpi.recvAntiMessages()" << endl;
    logFile.flush();
    mpi.recvAntiMessages(&antiMessages);

```

```

    // process waiting antimessages
    if(!antiMessages.empty()) {
//      logFile << "processing anti-messages!" << endl;
logFile.flush();
        // get the lowest time of any past antimessages
        for(vector<message>::iterator i=antiMessages.begin(); i <
antiMessages.end(); ++i) {
            if((*i).time <= localTime) {
                if((*i).time < pastTime || !isRollback) {
                    pastTime = (*i).time;
                    isRollback = true;
                }
            }
        }

        // if there is a past event rollback to the minimum past event
type
        if(isRollback) {
            // rollback to the time of this antimessage
            rollback(pastTime);
        }

        // translate the antimessages and insert them into the
antiEvent vector
        translateMessages(&antiEvents,&antiMessages);
    }

    // if we don't have any waiting messages just exit out
    if(messages.empty())
        return(true);

//  logFile << "processing messages!" << endl; logFile.flush();

    // translate the remote messages into events4
    translateMessages(&remoteEvents,&messages);

    // loop through all the recieved events and push them into the
isRollback = false; pastTime = 0.0;
    for(vector<Event*>::iterator i=remoteEvents.begin(); i < remoteEvents.end();
++i) {
        if((*i)->time < localTime) {
            if((*i)->time < pastTime || !isRollback) {
                pastTime = (*i)->time;
                isRollback = true;
            }
        }
    }
}

```

```

// if we need to rollback then do a rollback
if(isRollback)
    rollback(pastTime);

// push all the future events onto the remote event list
for(vector<Event*>::iterator i=remoteEvents.begin(); i < remoteEvents.end();
++i)
    remoteEventList.push(*i);

return(true);
}

bool Lattice::translateMessages(vector<Event*>* events, vector<message*>*
messages) {
    vector<message*>::iterator i;
    message m;

    for(i=messages->begin(); i < messages->end(); ++i) {
        m = *i;

        // make an event from it <sarcasm>this should be fun</sarcasm>
        switch(m.type) {
            case eventDiffusion: {
                // translate the oldSite and newSite coords
                if(m.newSite.p.x < LEFT_X_BOUNDARY) {
                    m.newSite.p.x = DIM_X - 1 - m.newSite.p.x;
                    m.oldSite.p.x = DIM_X - 1 - m.oldSite.p.x;
                } else {
                    m.newSite.p.x -= DIM_X - 1;
                    m.oldSite.p.x -= DIM_X - 1;
                }

                // push an event into the return vector
                events->push_back(new Event(&lattice[m.oldSite.p.x][m.oldSite.p.y],
&lattice[m.newSite.p.x][m.newSite.p.y],m.time,false,m.type,0));
            } break;
            case eventDeposition: {
                if(m.newSite.p.x < LEFT_X_BOUNDARY)
                    m.newSite.p.x = DIM_X - 1 - m.newSite.p.x;
                else
                    m.newSite.p.x -= DIM_X - 1;

                // push an event into the return vector
                events->push_back(new
Event(&lattice[m.newSite.p.x][m.newSite.p.y],m.time,false,m.type));
            } break;
            default:throw(Exception("Bad event type encountered in
Lattice::translateMessages()"));

```

```

    }
}

return(true);
}

message* Lattice::makeMessage(Event* event) {

    switch(event->eventType) {
        case eventDiffusion: {
            m.oldSite.p.x = event->oldSite->p.x;
            m.oldSite.p.y = event->oldSite->p.y;
            m.oldSite.h = event->oldSite->h;
            //m.oldSite = *(event->oldSite);
        } // fall through
        case eventDeposition: {
            m.newSite.p.x = event->newSite->p.x;
            m.newSite.p.y = event->newSite->p.y;
            m.newSite.h = event->newSite->h;
            //m.newSite = *(event->newSite);
        }break;
    }
    m.time = event->time;
    m.type = event->eventType;

    return(&m);
}

bool Lattice::hasAntiEvent(Event* event) {

    // if there are no antiEvents just bail
    if(antiEvents.empty())
        return(false);

    // loop through the available antiEvents
    for(vector<Event*>::iterator i=antiEvents.begin(); i < antiEvents.end();
    ++i) {
        // if the time of the event is the same as the time of the
        antievent
        // we can return true and erase the antievent
        if(event->time == (*i)->time) {
            delete *i; // clean up the allocated memory
            antiEvents.erase(i);
            return(true);
        }
    }

    // no matching antievent was found to return false
    return(false);
}

```

}

B.6 event.h

Code: event.h

```
#include "latprim.h"

#ifndef EVENT_H
#define EVENT_H

enum EventType {eventDeposition, eventDiffusion};

class Event {
public:
    Event(site* oldSite, site* newSite, double time, bool isLocal, EventType
eventType,/* Direction dir,*/ int listIndex) {
        this->oldSite = oldSite;
        this->newSite = newSite;
        this->time = time;
        this->isLocal = isLocal;
        this->eventType = eventType;
        /*this->dir = dir;*/
        this->listIndex = listIndex;
    }

    Event(site* newSite, double time, bool isLocal, EventType eventType/*,
Direction dir, int listIndex*/) {
        this->newSite = newSite;
        this->time = time;
        this->isLocal = isLocal;
        this->eventType = eventType;
        /*this->dir = dir;*/
        /*this->listIndex = listIndex;*/
    }

    site* oldSite;        // the old site (diffusion event)
    site* newSite;        // the new site (diffusion event) or the site
for a deposition event

```

```

    double time;           // event time
    bool isLocal;          // true if a local event
    bool isBoundry;        // true if a boundry event
    EventType eventType;   // type of event
    // Direction dir;      // important if it's a boundry event,
    // which neighbor to send it too
    int listIndex;         // index of the site in the list
private:
    Event() { ; } // supress creating an event with a default
    constructor
};

#endif

```

B.7 rewindlist.h

Code: rewindlist.h

```

#include <vector>
using std::vector;

#include <stack>
using std::stack;

#include "exception.h"

#ifndef REWINDLIST_H
#define REWINDLIST_H

enum EntryType {entryAdd, entryRemove};

template <class T>
class RewindEntry {
public:
    RewindEntry() { ; }
    RewindEntry(double t, T value, int oldPos, EntryType entryType) {
        this->t = t;
        this->value = value;
        this->oldPos = oldPos;
        this->entryType = entryType;
    }
};

```



```

    }

    // a collection of useful operator overloads (most could be
    removed)
    bool operator == (double right) {
        return(t == right);
    }

    bool operator >= (double right) {
        return(t <= right);
    }

    bool operator <= (double right) {
        return(t <= right);
    }

    bool operator < (double right) {
        return(t < right);
    }

    bool operator > (double right) {
        return(t > right);
    }

    T value;
    double t;
    int oldPos;
    EntryType entryType;
private:
};

template <class T>
class RewindList {
public:
    RewindList() { ; }
    RewindList(T*,int);

    bool rewind(double);
    bool stepBack();

    int add(T,double);
    T remove(int,double);

    T& operator [] (int i) { return(list[i]); }

    int size() { list.size(); }

    typedef vector<T>::iterator iterator;

```

```

private:
    vector<T> list;
    stack< RewindEntry<T> > rewindStack;
};

template <class T>
RewindList<T>::RewindList(T* ary, int size) {
    // seed the list with initial values
    for(int i = 0; i < size; ++i) {
        list.push_back(ary[i]);
    }
}

template <class T>
int RewindList<T>::add(T n, double t) {
    // push a new element on the back of the array
    list.push_back(n);

    // insert an event into our rewind stack so we can delete this
    // during a rollback
    rewindStack.push(RewindEntry<T>(t,n,0,entryAdd));

    // return the index of the last element of the array
    return(list.size() - 1);
}

template <class T>
T RewindList<T>::remove(int pos, double t) {
    RewindEntry<T> old(t,list[pos],pos,entryRemove);
    iterator index = list.begin() + pos;

    // push an event onto the rewind stack so we can rewind this
    // during a rollback
    rewindStack.push(old);

    if(list.size() <= 0 || index > list.end())
        return(NULL);

    // erase the item currently there
    list.erase(index);
    // insert the last item in the list there
    list.insert(index,list.back());
    // remove the last item from the list to remove the dupe
    list.pop_back();

    return(*index);
}

template <class T>

```

```

bool RewindList<T>::rewind(double t) {
    RewindEntry<T> top;
    iterator i;

    // loop until the rewind stack is empty or we've reached the
first event
    // before time t
    while(!rewindStack.empty() && rewindStack.top().t > t) {
        stepBack();
    }

    return(true);
}

template <class T>
bool RewindList<T>::stepBack() {
    RewindEntry<T> top;
    iterator i;
    double t;

    if(rewindStack.empty())
        return(false);

    // retrieve the time on the top of the stack
    t = rewindStack.top().t;

    // rewind the event
    while(t == rewindStack.top().t) {
        // retrieve the top of the stack
        top = rewindStack.top();

        switch(top.entryType) {
            case entryAdd: {
                // added an item to the end of the list so just pop it off
                list.pop_back();
            }break;
            case entryRemove: {
                // calculate the position to insert the old record
                i = list.begin() + top.oldPos;

                // push the item to the back of the list
                list.push_back(list[top.oldPos]);

                // erase the item
                list.erase(i);

                // restore the previous item
                list.insert(i,top.value);
            }break;
        }
    }
}

```

```
    }

    // pop the top of the rewind stack
    rewindStack.pop();

}

// return something hopeful
return(true);
}

#endif
```

B.8 randgen.h

Code: randgen.h

```
#include <vector>
using std::vector;

#include <stack>
using std::stack;

#include <iostream>
using std::cout;
using std::endl;

#ifndef RANDGEN_H
#define RANDGEN_H

class RandGen {
public:
    RandGen(int size) : expand(size) {
        this->seed = 0;
        populateList(size);
        position = randList.begin();
    }

    RandGen(int size,int seed) : expand(size) {
```

```

        this->seed = seed;
        populateList(size);
        position = randList.begin();
    }

    bool rewind(double t) {
        if(times.empty())
            return(true);

        while(!times.empty() && times.top() > t) {
            position--;
            times.pop();
        }

        if(position <= randList.begin())
            position = randList.begin();
        return(true);
    }

    float getRandom(double t) {
        ++position;
        times.push(t);
        if(position != randList.end())
            return(*(position));
        else {
            int offset = randList.size();
            populateList(expand);
            position = randList.begin() + offset;
        }
        return(*(position));
    }
}

private:
    vector<float> randList;
    vector<float>::iterator position;
    stack<double> times;

    int seed;
    int expand;

    RandGen() { ; }

    void populateList(int count) {
        for(int i=0; i < count; ++i)
            randList.push_back((float)genRand());
    }

    double genRand() {

```

```

/*-----*/
/* A C-program for TT800 : July 8th 1996 Version */
/* by M. Matsumoto, email: matumoto@math.keio.ac.jp */
/* genrand() generate one pseudorandom number with double
precision */
/* which is uniformly distributed on [0,1]-interval */
/* for each call. One may choose any initial 25 seeds */
/* except all zeros. */

/* See: ACM Transactions on Modelling and Computer Simulation,
*/
/* Vol. 4, No. 3, 1994, pages 254-266. */

const int NRan = 25;
int MRan = seed % NRan;

unsigned long y;
static int k = 0;
static unsigned long x[NRan]={ /* initial 25 seeds, change as you
wish */
                                0x95f24dab, 0x0b685215, 0xe76ccae7,
0xaf3ec239, 0x715fad23,
                                0x24a590ad, 0x69e4b5ef, 0xbf456141,
0x96bcb7b, 0xa7bdf825,
                                0xc1de75b7, 0x8858a9c9, 0x2da87693,
0xb657f9dd, 0xffdc8a9f,
                                0x8121da71, 0x8b823ecb, 0x885d05f5,
0x4e20cd47, 0x5a9ad5d9,
                                0x512c0c03, 0xea857ccd, 0x4cc1d30f,
0x8891a8a1, 0xa6b7aadb
                                };
static unsigned long mag01[2]={0x0, 0x8ebfd028 /* this is magic vector
'a', don't change */};
if (k==NRan) { /* generate NRan words at one time */
    int kk;
    for (kk=0;kk<NRan-MRan;kk++) {
        x[kk] = x[kk+MRan] ^ (x[kk] >> 1) ^ mag01[x[kk] % 2];
    }
    for (; kk<NRan;kk++) {
        x[kk] = x[kk+(MRan-NRan)] ^ (x[kk] >> 1) ^ mag01[x[kk] % 2];
    }
    k=0;
}
y = x[k];
y ^ = (y << 7) & 0x2b5b2500; /* s and b, magic vectors */
y ^ = (y << 15) & 0xdb8b0000; /* t and c, magic vectors */
y &= 0xffffffff; /* you may delete this line if word size = 32 */
/*
    the following line was added by Makoto Matsumoto in the

```

```

1996 version
    to improve lower bit's corellation.
    Delete this line to o use the code published in 1994.
    */
    y ^ = (y >> 16); /* added to the 1994 version */
    k++;
    return( (double) y / (unsigned long) 0xffffffff);
}

};

#endif

```

B.9 mpiwrapper.h

Code: mpiwrapper.h

```

#include <vector>
using std::vector;

#include <fstream>
using std::fstream;
using std::endl;

#include <iostream>
using std::cout;

#include <mpi.h>

#include "latprim.h"
#include "latconst.h"
#include "exception.h"
#include "event.h"

#ifndef MPIWRAPPER_H
#define MPIWRAPPER_H

#define LEFT(a,b) (((a - 1) >= 0) ? (a - 1) : -1)

```

```

#define RIGHT(a,b) (((a + 1) < b) ? (a + 1) : -1)
#define DIR(a) ((a == LEFT) ? left : right)
const int BUFFER_SIZE_COUNT = 1024*1024*10; // 10MB buffer (overkill,
baby)
const int ROOT_RANK = 0;
const int NUM_NEIGHBORS = 2;

const int TAG_MESSAGE = 0;
const int TAG_ANTI_MESSAGE = 1;

enum Direction {LEFT,RIGHT};

typedef struct {
    site oldSite;
    site newSite;
    double time;
    EventType type;
} message;

class MPIWrapper {
public:
    MPIWrapper();
    ~ MPIWrapper();

    bool init(int*, char***);
    bool shutdown();

    bool sendMessage(message* , Direction);
    bool recvMessages(vector<message>*);
    bool isMessage();

    bool sendAntiMessage(message* , Direction);
    bool recvAntiMessages(vector<message>*);
    bool isAntiMessage();

    float allReduceFloat(float, MPI_Op);
    double allReduceDouble(double, MPI_Op);
    int allReduceInt(int, MPI_Op);

    void barrier();
    double wallTime();

    bool isRoot();

    int getRank() {
        return(rank);
    }

    int getNodeCount() {

```



```

    return(nodeCount);
}

void printStats(fstream& file) {
    file << "--- MPI STATS ---" << endl;
    file << "rank = " << rank << endl;
    file << "nodeCount = " << nodeCount << endl;
    file << "left = " << left << " right = " << right << endl;
    file << "Send Messages = " << countSend << " Recieved Messages = " <<
countRecv << endl;
    file << "Send Anti-Messages = " << countSendAnti << " Recieved
Anti-Messages = " << countRecvAnti << endl;
    file << "Total Send = " << (countSend + countSendAnti) << " Total Recv = "
<< (countRecv + countRecvAnti) << endl;
    file << "-----" << endl << endl;
    file.flush();
}

private:
    int rank;
    int nodeCount;
    bool isInit;

    char* buffer;
    int bufferSize;

    MPI_Datatype typeSite;
    MPI_Datatype typePoint;
    MPI_Datatype typeMessage;

    int left;
    int right;
    message m;

    MPI_Status status;
    int flag;

    int countSend;
    int countRecv;
    int countSendAnti;
    int countRecvAnti;
};

#endif

```

B.10 mpiwrapper.cpp

Code: mpiwrapper.cpp

```
#include <vector>
using std::vector;

#include <mpi.h>
#include <stdlib.h>

#include "mpiwrapper.h"
#include "latprim.h"
#include "exception.h"
#include "event.h"

MPIWrapper::MPIWrapper() : rank(-1), nodeCount(-1), isInit(false), left(-1),
right(-1), countSend(0), countRecv(0), countSendAnti(0), countRecvAnti(0) { ; }

MPIWrapper::~MPIWrapper() {
    //if(isInit)
    //    shutdown();
}

bool MPIWrapper::init(int* argv, char** argc[]) {
    MPI_Aint* displacements;
    MPI_Datatype* dataTypes;
    int* blockLength;
    MPI_Aint startAddress;
    MPI_Aint address;
    point p;
    site s;

    // see if init == true, if that is so we've got big problems
    if(isInit)
        throw(Exception("ERROR: Duplicate call to MPIWrapper::init(!)"));

    // call MPI_Init() to start this whole shebang
    MPI_Init(argv, argc);

    // get the process rank and the number of nodes
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nodeCount);
```

```

// make sure the shit didn't hit the fan
if(rank < 0) {
    throw(Exception("ERROR: MPI_Comm_rank() failed to return useful value!"));
}
if(nodeCount < 0) {
    throw(Exception("ERROR: MPI_Comm_size() failed to return useful value!"));
}

// create the datatype for the point structure
displacements = new MPI_Aint[2];
dataTypes = new MPI_Datatype[2];
blockLength = new int[2];

blockLength[0] = 1;
blockLength[1] = 1;
dataTypes[0] = MPI_INT;
dataTypes[1] = MPI_INT;

MPI_Address(&p.x, &startAddress);
displacements[0] = 0;
MPI_Address(&p.y, &address);
displacements[1] = address - startAddress;

MPI_Type_struct(2, blockLength, displacements, dataTypes, &typePoint);
MPI_Type_commit(&typePoint);

delete [] displacements;
delete [] dataTypes;
delete [] blockLength;

// create the datatype for the site structure
displacements = new MPI_Aint[3];
dataTypes = new MPI_Datatype[3];
blockLength = new int[3];

blockLength[0] = 1;
blockLength[1] = 1;
blockLength[2] = 1;
dataTypes[0] = typePoint;
dataTypes[1] = MPI_INT;
dataTypes[2] = MPI_INT;

MPI_Address(&s.p, &startAddress);
displacements[0] = 0;
MPI_Address(&s.listIndex, &address);
displacements[1] = address - startAddress;
MPI_Address(&s.h, &address);
displacements[2] = address - startAddress;

```

```

MPI_Type_struct(3,blockLength,displacements,dataTypes,&typeSite);
MPI_Type_commit(&typeSite);

delete [] displacements;
delete [] dataTypes;
delete [] blockLength;

// create the datatype for the boundryEvent structure
displacements = new MPI_Aint[4];
dataTypes = new MPI_Datatype[4];
blockLength = new int[4];

blockLength[0] = 1;
blockLength[1] = 1;
blockLength[2] = 1;
blockLength[3] = 1;
dataTypes[0] = typeSite;
dataTypes[1] = typeSite;
dataTypes[2] = MPI_DOUBLE;
dataTypes[3] = MPI_INT;

MPI_Address(&m.oldSite,&startAddress);
displacements[0] = 0;
MPI_Address(&m.newSite,&address);
displacements[1] = address - startAddress;
MPI_Address(&m.time,&address);
displacements[2] = address - startAddress;
MPI_Address(&m.type,&address);
displacements[3] = address - startAddress;

MPI_Type_struct(4,blockLength,displacements,dataTypes,&typeMessage);
MPI_Type_commit(&typeMessage);

delete [] displacements;
delete [] dataTypes;
delete [] blockLength;

// attach the buffer to the MPI process
MPI_Buffer_attach(malloc(BUFFER_SIZE_COUNT * sizeof(message) +
MPI_BSEND_OVERHEAD), BUFFER_SIZE_COUNT * sizeof(message) +
MPI_BSEND_OVERHEAD);

// get the node on my left
left = LEFT(rank,nodeCount);

// get the node on my right
right = RIGHT(rank,nodeCount);

// hey, we finished the init! so set the flag

```

```

    isInit = true;

    // return the value of the flag (should be true)
    return(isInit);
}

bool MPIWrapper::shutdown() {
    // make sure we had a successful init() call
    if(!isInit)
        return(false);

    // detach the buffer from the MPI process (COULD STALL PROGRAM
    EXECUTION
    // SINCE ALL BUFFERED MESSAGES MUST BE DELIVERED BEFORE THE CALL
    EXITS)
    MPI_Buffer_detach(&buffer, &bufferSize);

    // free the declared types
    MPI_Type_free(&typeMessage);
    MPI_Type_free(&typeSite);
    MPI_Type_free(&typePoint);

    // call the MPI_Finalize() function to make MPI clean up
    MPI_Finalize();

    // set init to false so we don't do anything stupid
    isInit = false;

    // return true so all is well
    return(!isInit);
}

bool MPIWrapper::sendMessage(message* m, Direction dir) {

    // send the message with a buffered send so we don't block
    if(DIR(dir) != -1) {
        MPI_Bsend(m, 1, typeMessage, DIR(dir), TAG_MESSAGE, MPI_COMM_WORLD);
        ++countSend;
    }

    // return true
    return(true);
}

bool MPIWrapper::recvMessages(vector<message>* messages) {

    // loop until we don't have any more messages waiting
    while(isMessage()) {
        // recieve the message

```

```

        MPI_Recv(&m, 1, typeMessage, MPI_ANY_SOURCE, TAG_MESSAGE, MPI_COMM_WORLD,
&status);
        messages->push_back(m);
        ++countRecv;
    }

    // return true
    return(true);
}

bool MPIWrapper::isMessage() {
    // do an iprobe to get the value of flag (TRUE OR FALSE)
    MPI_Iprobe(MPI_ANY_SOURCE, TAG_MESSAGE, MPI_COMM_WORLD, &flag, &status);

    // return the value compared to the true equiv of 1 (since it's
an int)
    return(flag == 1);
}

bool MPIWrapper::sendAntiMessage(message* m, Direction dir) {

    // send the message with a buffered send so we don't block
    if(DIR(dir) != -1) {
        MPI_Bsend(m, 1, typeMessage, DIR(dir), TAG_ANTI_MESSAGE, MPI_COMM_WORLD);
        ++countSendAnti;
    }

    // return true
    return(true);
}

bool MPIWrapper::recvAntiMessages(vector<message>* messages) {

    // loop until we don't have any more messages waiting
    while(isAntiMessage()) {
        // recieve the message
        MPI_Recv(&m, 1, typeMessage, MPI_ANY_SOURCE, TAG_ANTI_MESSAGE,
MPI_COMM_WORLD, &status);
        messages->push_back(m);
        ++countRecvAnti;
    }

    // return true
    return(true);
}

bool MPIWrapper::isAntiMessage() {
    // do an iprobe to get the value of flag (TRUE OR FALSE)
    MPI_Iprobe(MPI_ANY_SOURCE, TAG_ANTI_MESSAGE, MPI_COMM_WORLD, &flag,

```

```

&status);

    // return the value compared to the true equiv of 1 (since it's
    an int)
    return(flag == 1);
}

float MPIWrapper::allReduceFloat(float input, MPI_Op op) {
    float output;

    // call MPI_Allreduce() using the provided input/output, the
    correct datatype
    // and the user-provided op for the world communicator
    MPI_Allreduce(&input,&output,1,MPI_FLOAT,op,MPI_COMM_WORLD);

    // return the output value
    return(output);
}

double MPIWrapper::allReduceDouble(double input, MPI_Op op) {
    double output;

    // call MPI_Allreduce() using the provided input/output, the
    correct datatype
    // and the user-provided op for the world communicator
    MPI_Allreduce(&input,&output,1,MPI_DOUBLE,op,MPI_COMM_WORLD);

    // return the output value
    return(output);
}

int MPIWrapper::allReduceInt(int input, MPI_Op op) {
    int output;

    // call MPI_Allreduce() using the provided input/output, the
    correct datatype
    // and the user-provided op for the world communicator
    MPI_Allreduce(&input,&output,1,MPI_INT,op,MPI_COMM_WORLD);

    // return the output value
    return(output);
}

bool MPIWrapper::isRoot() {
    // return the value of this compare
    return(rank == ROOT_RANK);
}

void MPIWrapper::barrier() {

```

```
MPI_Barrier(MPI_COMM_WORLD);  
}  
  
double MPIWrapper::wallTime() {  
    return(MPI_Wtime());  
}
```

B.11 exception.h

Code: exception.h

```
#include <string>  
using std::string;  
  
#ifndef EXCEPTION_H  
#define EXCEPTION_H  
  
class Exception {  
public:  
    Exception(char* c) : error(c) { ; }  
    Exception(string &s) : error(s) { ; }  
    Exception() : error("General Error") { ; }  
  
    string error;  
  
};  
  
#endif
```

Bibliography

- [1] Orna Berry and Greg Lomow. Speeding up distributed simulation using the time warp mechanism. In *EW 2: Proceedings of the 2nd workshop on Making distributed systems work*, pages 1–14, New York, NY, USA, 1986. ACM Press.
- [2] Paul Coddington. Monte carlo simulation for statistical physics. http://www.npac.syr.edu/users/paulc/lectures/montecarlo/p_montecarlo.html.
- [3] A. Esteve. Kinetic monte carlo technique in the modeling of micro/nanotechnologies. <http://www.laas.fr/ATOMCAD/Doc/aesteve.pdf>.
- [4] Boris Lubachevsky and Alan Weiss. Synchronous relaxation for parallel ising spin simulations. In *PADS '01: Proceedings of the fifteenth workshop on Parallel and distributed simulation*, pages 185–192, Washington, DC, USA, 2001. IEEE Computer Society.
- [5] Yimsic Shim and Jaques G. Amar. Synchronous relaxation algorithm for parallel kinetic monte carlo. *Physical Review B*, 71:1–13, 2005.